

# Blockchains & Cryptocurrencies

## **Smart Contracts & Ethereum**



Instructor: Matthew Green  
Fall 2020

Some slides adapted from NBFMG

# Housekeeping

- A1 in, A2 out (maybe) Weds (maybe Fri)

News?

# Today

- From Bitcoin to smart contracts
- Ethereum
- Applications of smart contracts



# Review: Bitcoin

- **Each block is a list of transactions**

- Each transaction consumes one or more inputs;
- Each transaction includes a set of outputs (amount + destination)
- Input consumption has conditions (e.g., valid script, typically enforcing valid signature)

# Review: Bitcoin

- **Each block is a list of transactions**
  - Each transaction consumes one or more inputs;
  - Each transaction includes a set of outputs (amount + destination)
  - Input consumption has conditions (e.g., valid signature)

In practice, each input/output has script:

ScriptPubKey (outputs)

ScriptSig (inputs)

# Review: Bitcoin

- |   |            |                            |
|---|------------|----------------------------|
| 76  | A9         | 14                         |
| OP_DUP  | OP_HASH160 | Bytes to push              |
| 89 AB CD EF AB BA AB BA AB BA AB BA AB BA AB BA AB BA 88 AC |            |                            |
| Data to push  |            | OP_EQUALVERIFY OP_CHECKSIG |

- Each transaction includes a set of outputs

```
scriptPubKey: OP_2DUP OP_EQUAL OP_NOT OP_VERIFY OP_SHA1 OP_SWAP OP_SHA1 OP_EQUAL
scriptSig: <preimage1> <preimage2>
```

- Input consumption has conditions (e.g., valid signature)

```
scriptPubKey: <expiry time> OP_CHECKLOCKTIMEVERIFY OP_DROP OP_DUP OP_HASH160 <pubKeyHash> OP_EQUALVERIFY OP_CHECKSIG
scriptSig: <sig> <pubKey>
```

ScriptPubKey (outputs)  
ScriptSig (inputs)

# Review: Bitcoin script

- **Bitcoin script allows us to attach conditions to payments**
  - However script is deliberately *limited*
  - Stack-based FORTH-type language, many original opcodes disabled
  - Highly limited access to global data



# Limitations & improvements

# Hard-coded limits in Bitcoin

- 10 min. average creation time per block
- 1 M bytes in a block
- 20,000 signature operations per block
- 23M total bitcoins maximum
- 50,25,12.5... bitcoin mining reward

These affect  
economic  
balance of  
power too  
much to  
change now

# Throughput limits in Bitcoin

- 1 M bytes/block (10 min)
- >250 bytes/transaction
- 7 transactions/block

Improving throughput: A  
strong motivation for  
new cryptocurrencies/  
DLT systems

Comparison

- VISA: 2,000-10,000 transactions/sec
- PayPal: 50-100 transaction/sec

# Cryptographic limits in Bitcoin

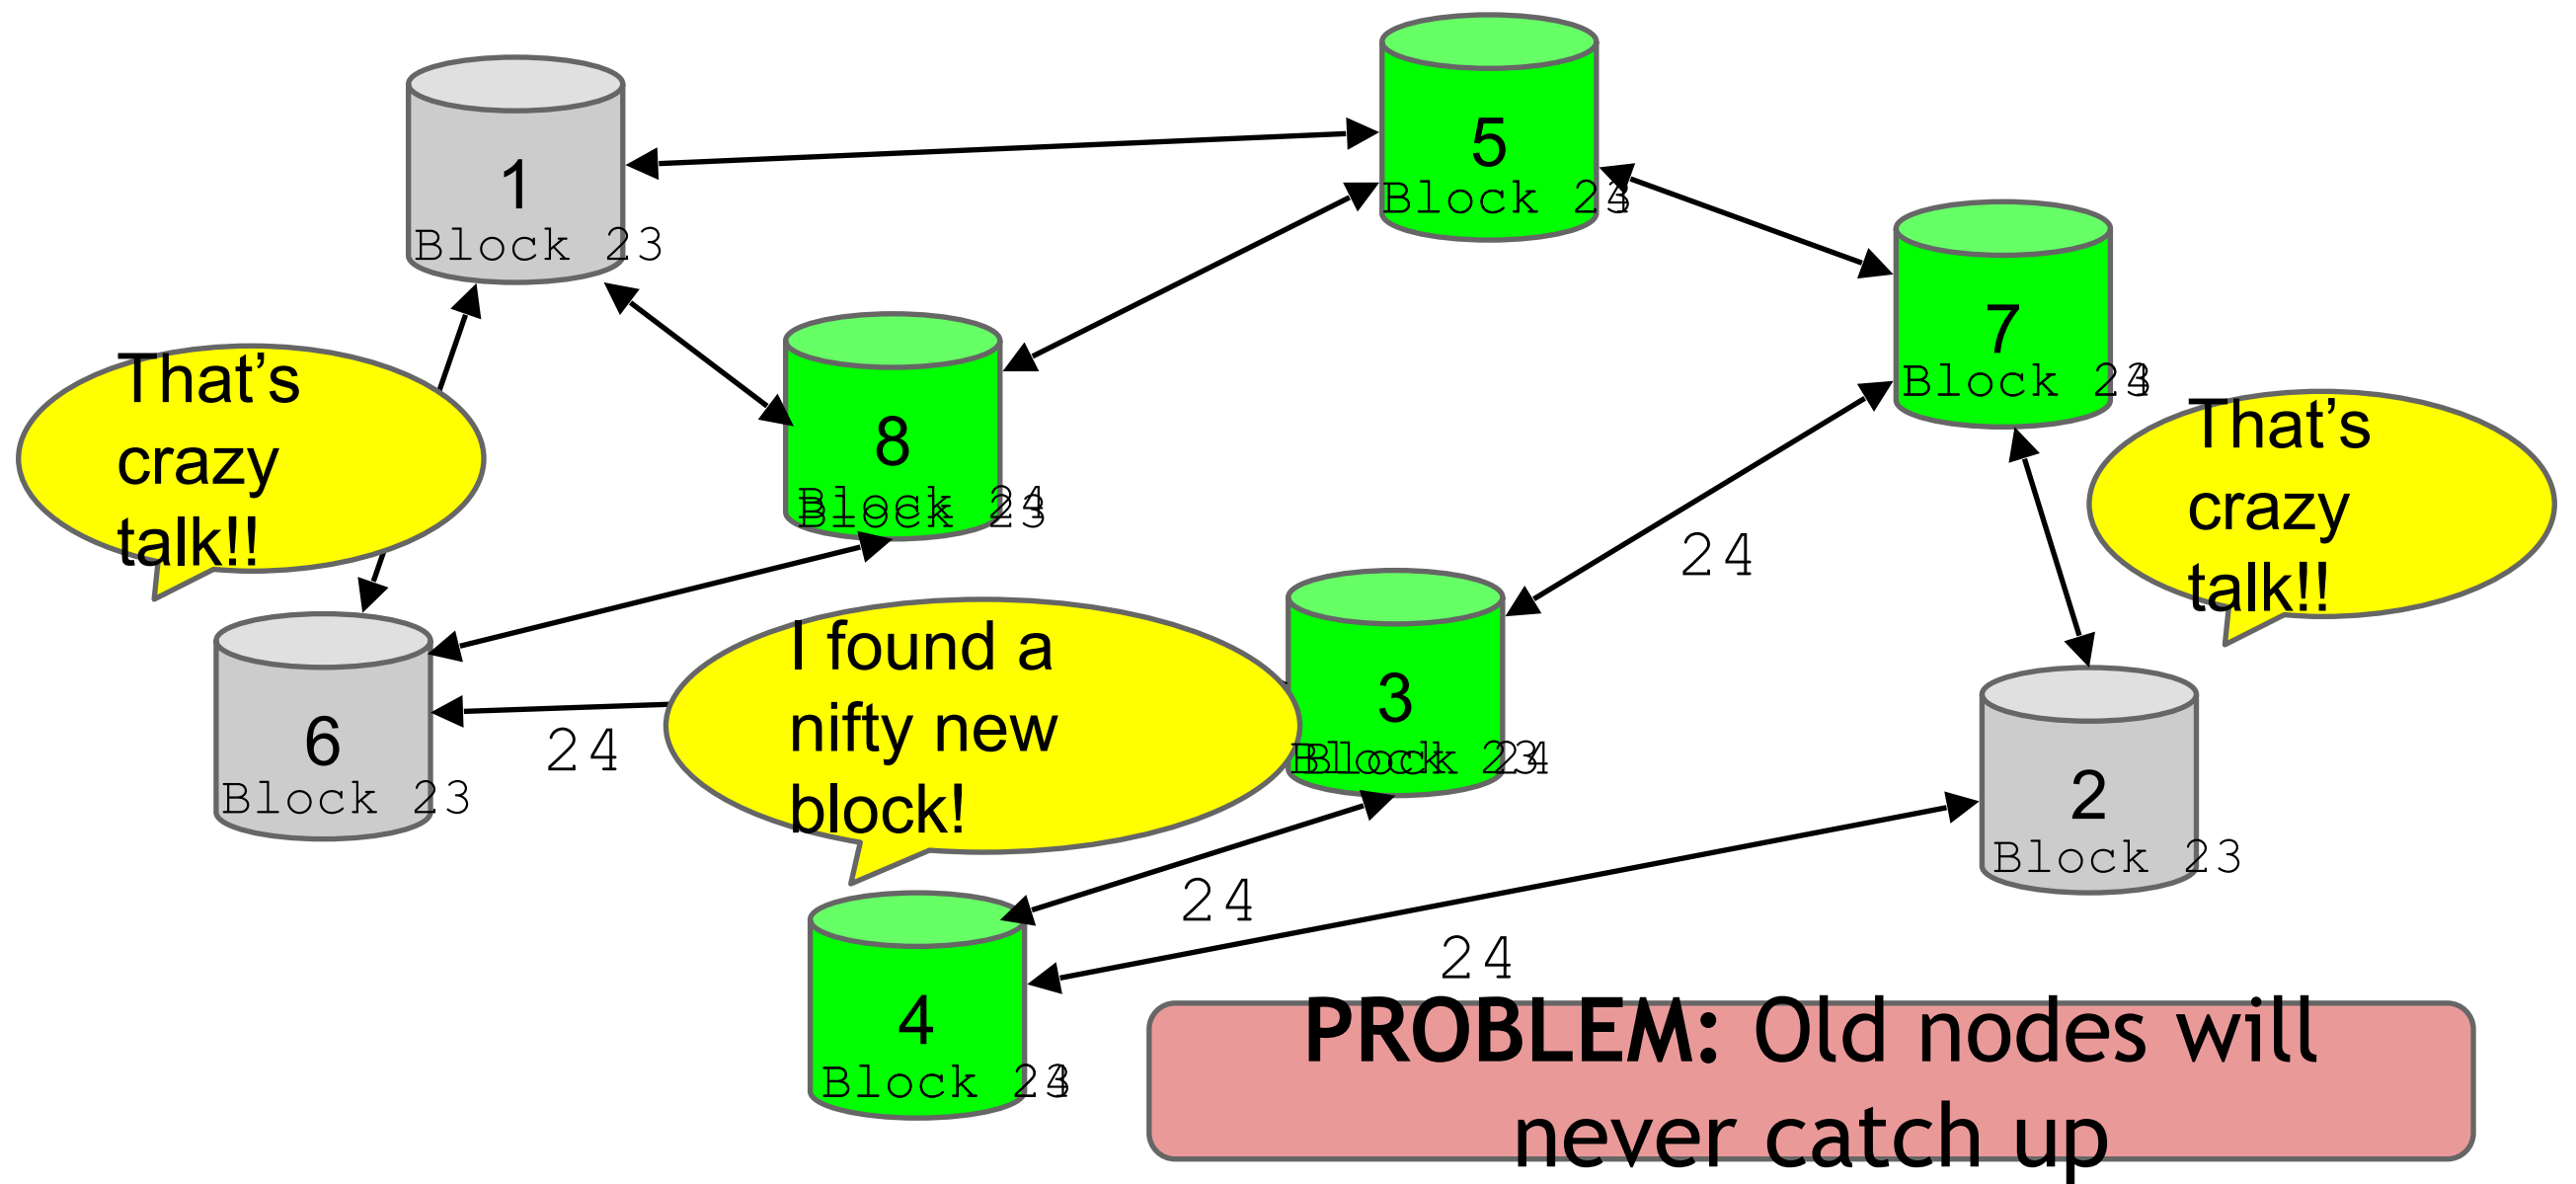
- Only 1 signature algorithm (ECDSA/P256)
- Schnorr signatures coming “soon”
- Hard-coded hash functions

Some of these crypto primitives used here might break by 2040 (e.g., collision-found in hash function, or powerful quantum computer breaks ECDSA)...

# Why not update Bitcoin software to overcome these limitations?

- Many of these changes require “hard forks”, which are currently considered unacceptable

# “Hard-forking” changes to Bitcoin



# Soft forks

Observation: we can add new features which only *limit* the set of valid transactions

Need majority of nodes to enforce new rules

Old nodes will approve

**RISK:** Old nodes might mine now-invalid blocks

# Soft fork example: pay to script hash

<signature>

<<pubkey> OP\_CHECKSIG>

OP\_HASH160

<hash of redemption script>

OP\_EQUAL

Old nodes will just approve the hash, not run the  
embedded script



# Soft fork possibilities

- New signature schemes
- Extra per-block metadata
  - Shove in the coinbase parameter
  - Commit to unspent transaction tree in each block

# Hard forks

- New op codes
- Changes to size limits
- Changes to mining rate
- Many small bug fixes

Currently seem unlikely to happen

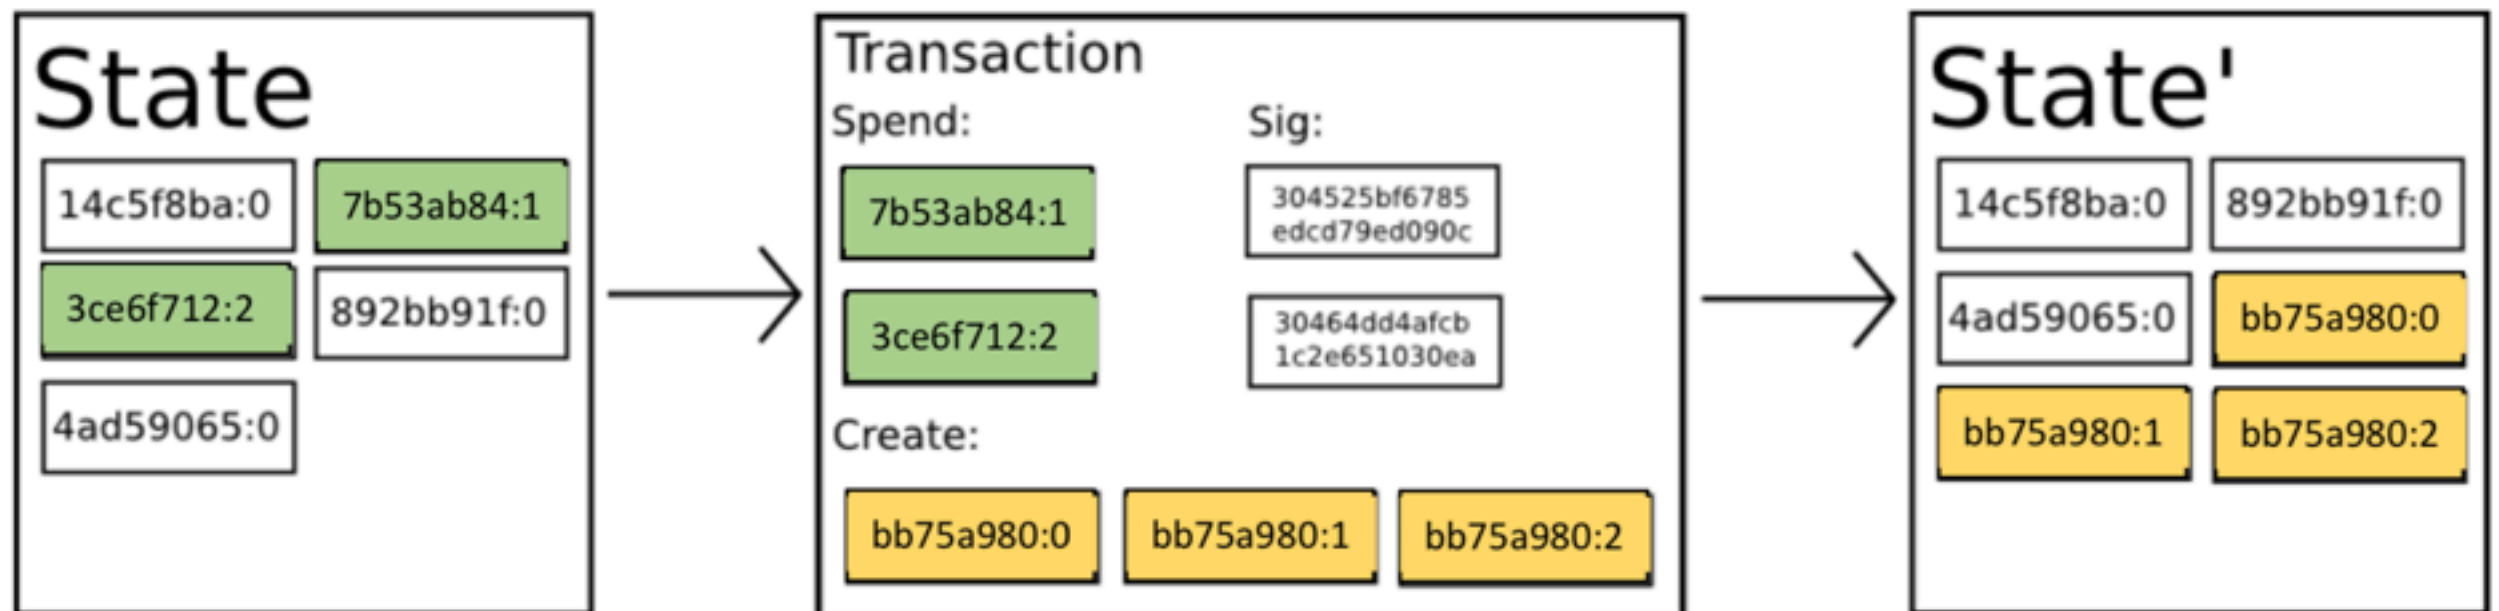
Many of these issues addressed by Altcoins

# Generalizing Bitcoin

- **Let's consider each Bitcoin transaction as a state transition function**
  - What is the input state?
  - What is the output state?
  - What does a Transaction do to the state?

# Generalizing Bitcoin






- **Let's consider each Bitcoin transaction as a state transition function**
  - Input state: list of coins available for spending (UTXO set)
  - Transaction: set of instructions for updating the UTXO set
  - Output state: New blockchain hash, Updated UTXO set



# “Smart contracts”

- Idea proposed by Nick Szabo (1994)
  - Bitcoin script is a ‘contract’ in the sense that it provides programmable conditions for redeeming a coin
  - However, conditions are highly limited
    - Example: pay out a coin iff a user has a signing key
    - Example: implement a second currency/asset
    - Example: pay out a coin iff a candidate wins the US election
    - Example: pay out a coin iff a majority of users votes to invest in a service

# “Smart contracts”

- Idea proposed by Nick Szabo (1994)
- Bitcoin script is a ‘contract’ in the sense that it provides programmable conditions for redeeming a coin
- However, conditions are highly limited
  - Example: pay out a coin iff a user has a signing key Bitcoin 
  - Example: implement a second currency/asset Bitcoin  
  - Example: pay out a coin iff a candidate wins the US Bitcoin 
  - Example: pay out a coin iff a majority of users votes to invest in a service Bitcoin 

## **Output coin condition:**

This coin can be paid out to user **A** on 1/20/2021 if the transaction contains a signature (under some “notary” public key) of the following candidate name (**X**)

Or

This coin can be paid out to user **B** on 1/20/2021 if the transaction contains a signature (under some “notary” public key) of the following candidate name (**Y**)

# Why not Bitcoin?

- Bitcoin script is not Turing complete
  - So script size (hence TX size) grows with complexity
- Bitcoin blocks are capped at 1 MB, which means TX space is at a premium
- Limited script opcodes (security reasons)
- No access to “global state”




# Ethereum

- Proposed in 2013 by Vitalik Buterin
  - Basic idea: extend Bitcoin by adding Turing-complete scripting, with full access to chain state (“smart contracts”)
  - Scripts run inside of an Ethereum virtual machine (EVM), can call other scripts & each other (recursively)
- Includes a native token (ETH) to pay for transactions, but users can create additional tokens using contracts

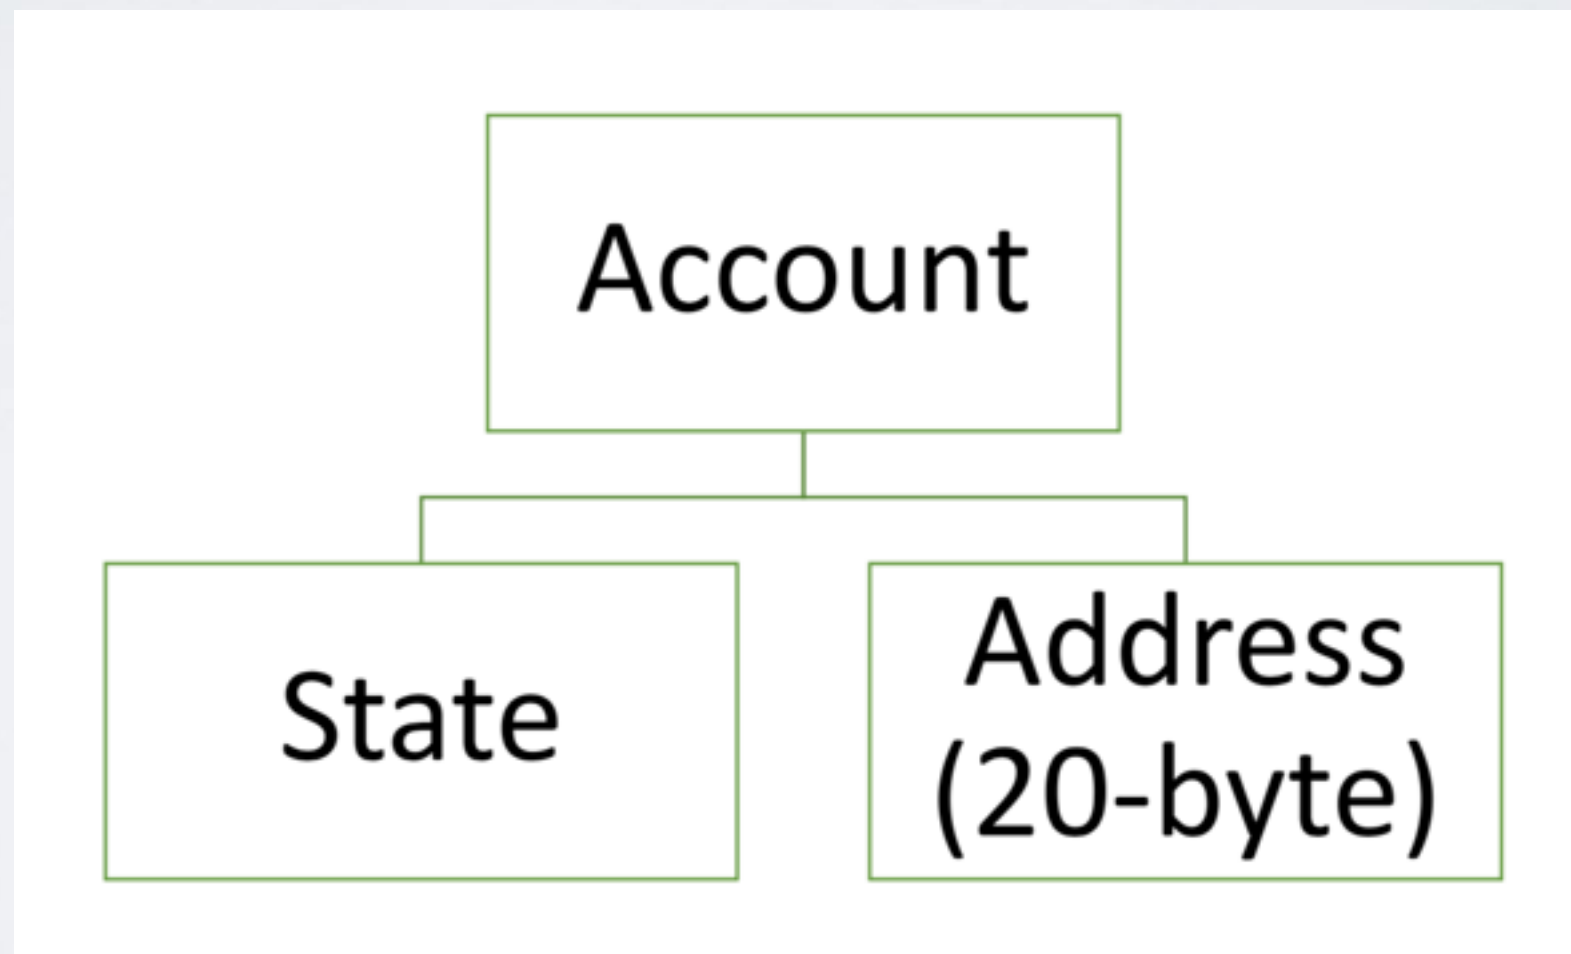


# What problems does this raise?



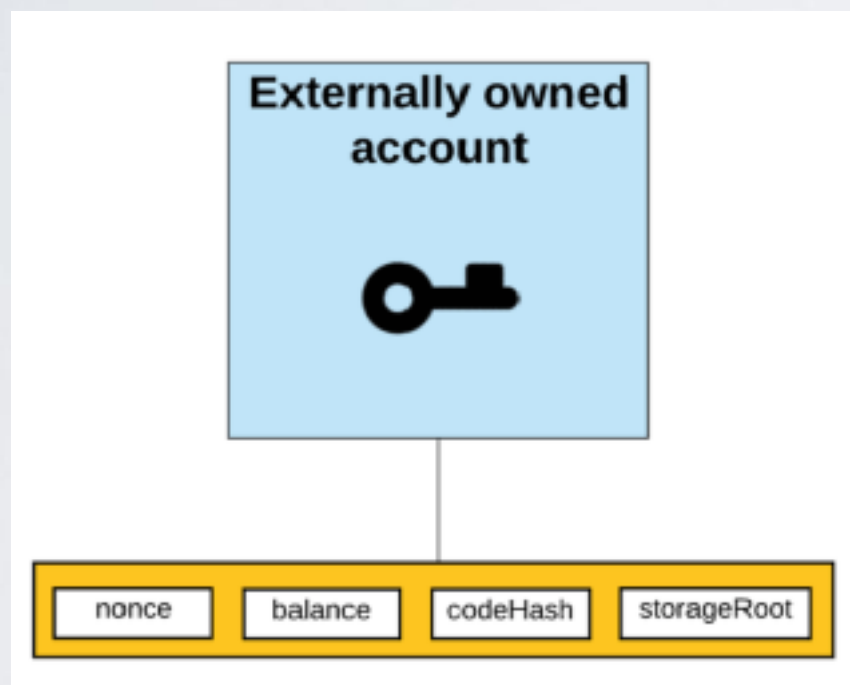
# Ethereum

- Account-based currency
  - Many small “accounts”

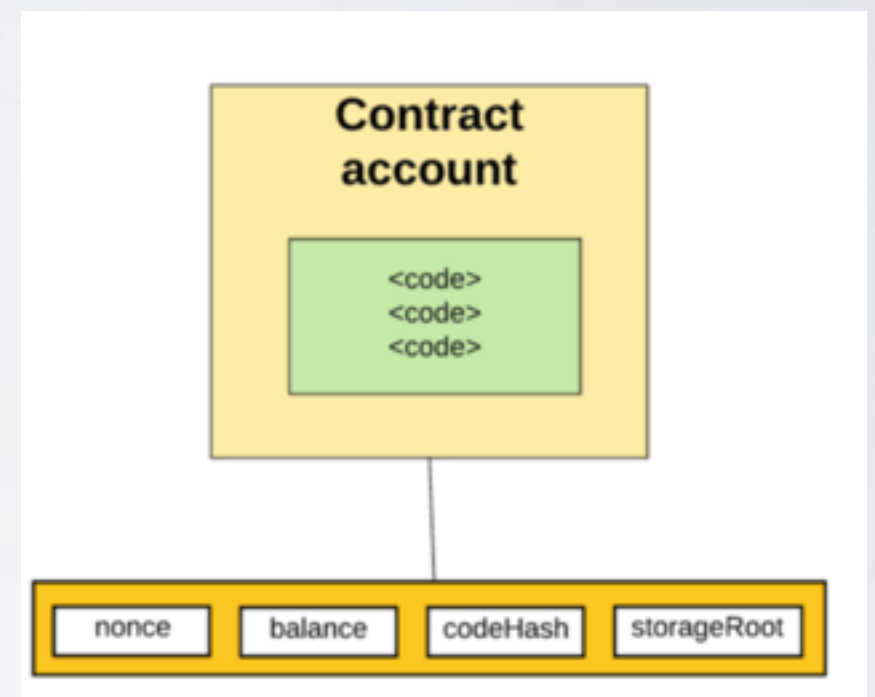


# Ethereum: Accounts

- Two types of account:
  - External (like Bitcoin), Contract accounts



Like Bitcoin, updates require an signature by an external private key



Anyone can call “methods” in the code, which trigger updates. Anyone can create.

# Ethereum: Accounts

**nonce:** # transactions sent/ # contracts created

- **balance:** # Wei owned ( $1 \text{ ether} = 10^{18} \text{ Wei}$ )

- **storageRoot:** Hash of the root node of a Merkle Patricia tree. The tree is empty by default.

- **codeHash:** Hash of empty string / Hash of the EVM (Ethereum Virtual Machine) code of this account

# Ethereum: Accounts

## Merkle Patricia tree:

A data structure for storing key/value pairs in a cryptographically authenticated manner.

(I.e., the tree root is a hash of all key/values in the structure, and updates/deletions are fast.)

nonce

- balance

- storage

tree is

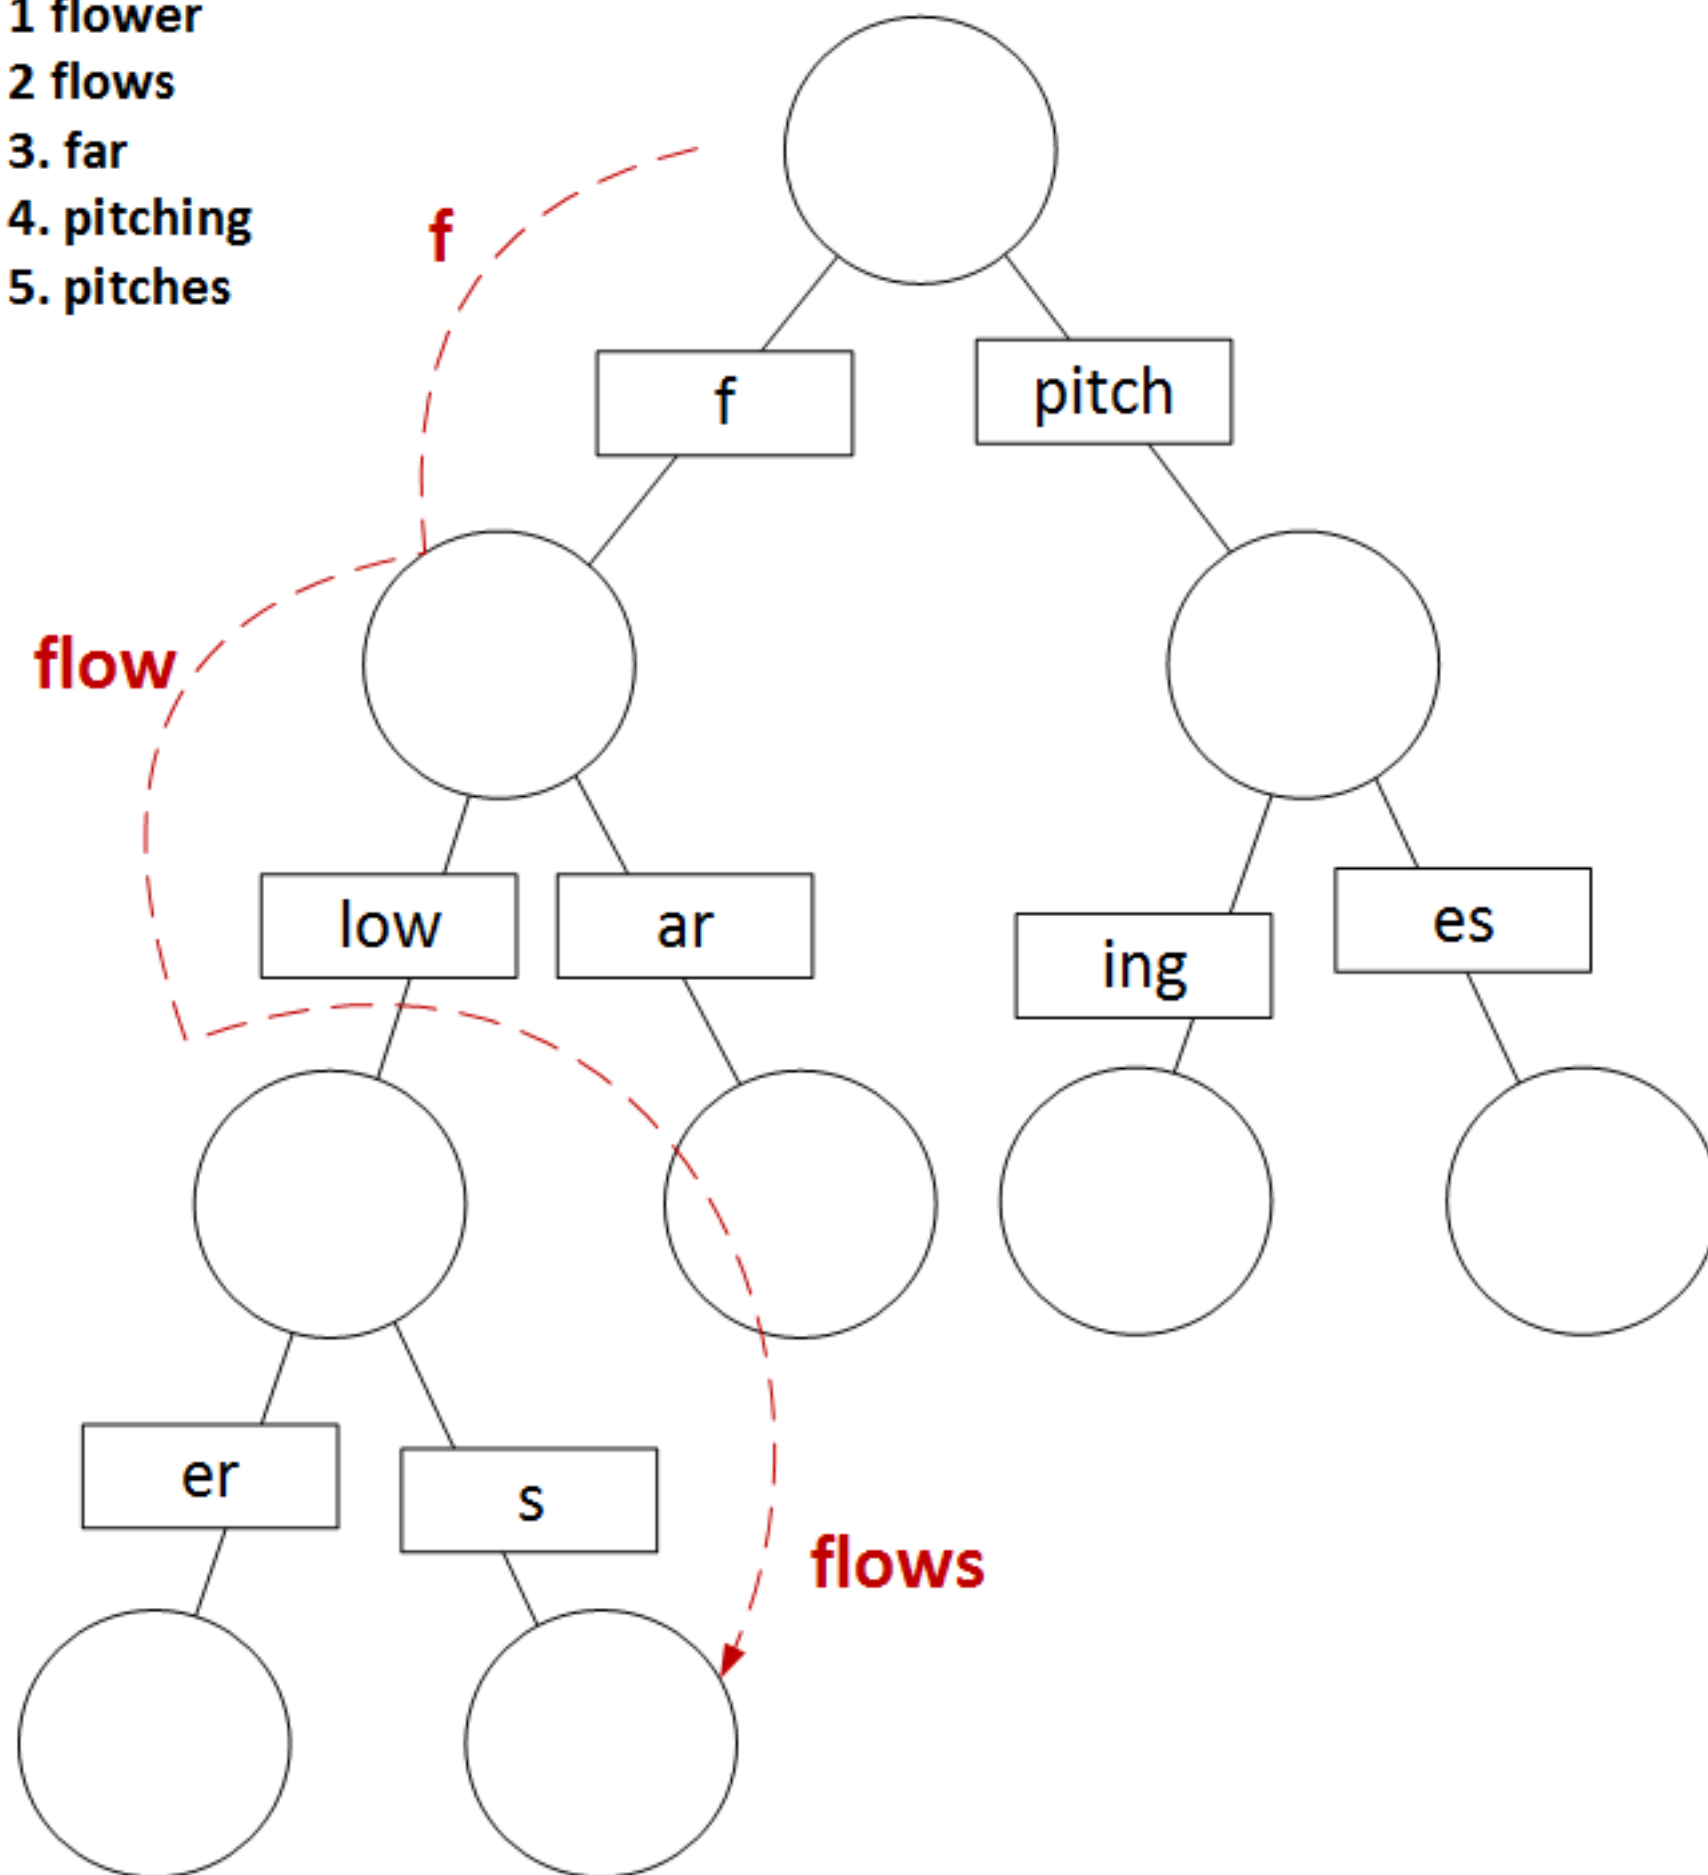
- code

Virtual

e. The

um

- 1 flower
- 2 flows
- 3. far
- 4. pitching
- 5. pitches



nonce

• bal

• stor

tree is

• cod

Virtu

(l.e., t  
st


e. The

eum

the

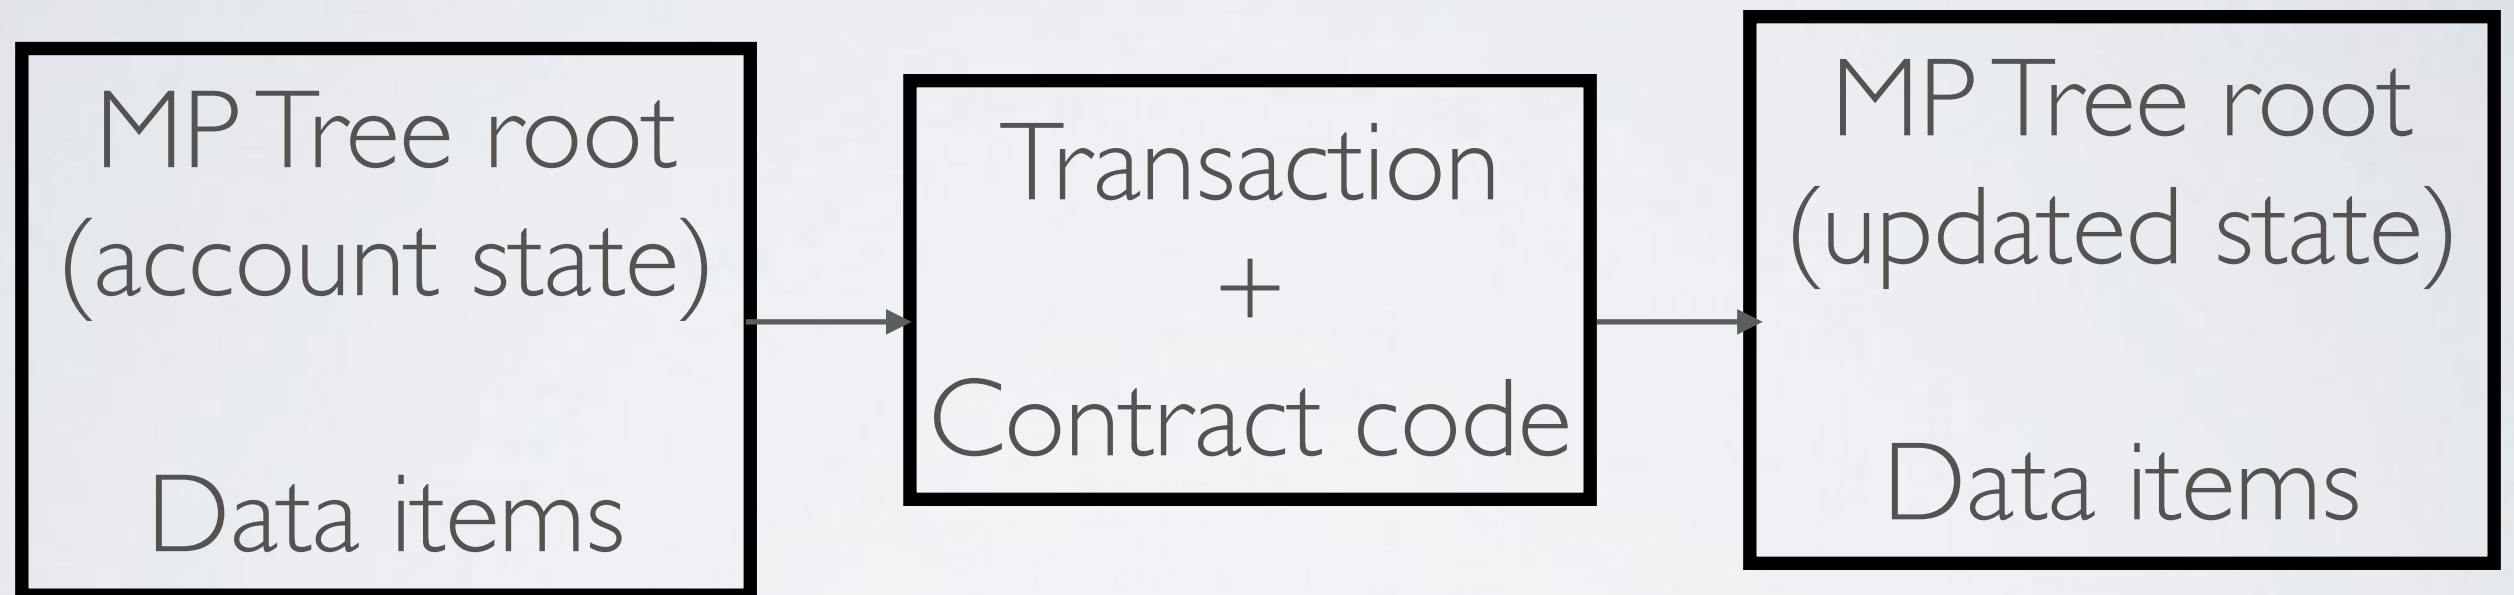
# Ethereum Transactions

- Two types:
  - Message calls: update state in a given contract, by executing code (or simply transferring money)
  - Contract creations: make a new contract account, with new state

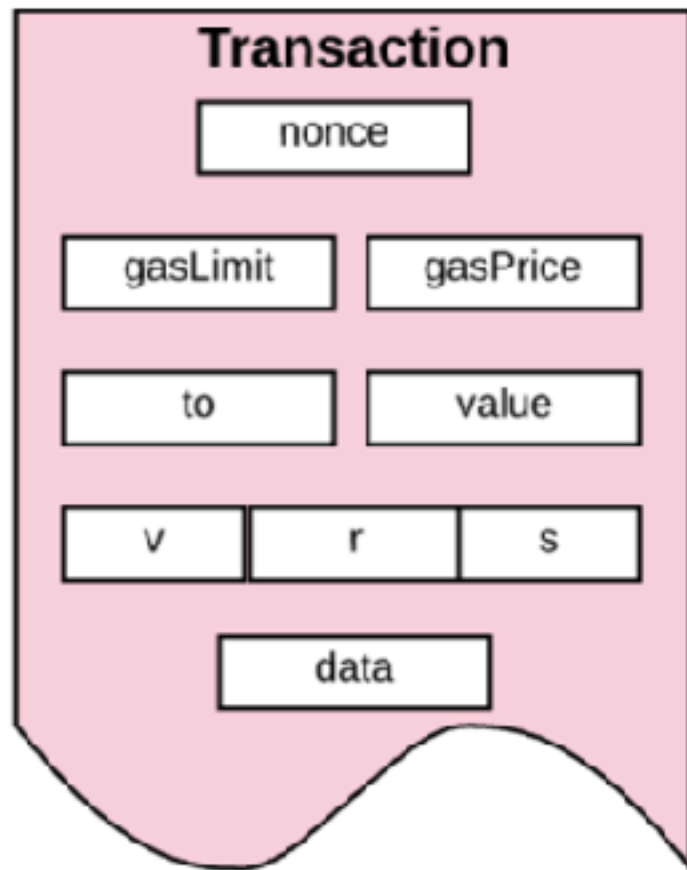




# Ethereum: Accounts



# Ethereum Transactions



- **nonce**: A count of the number of transactions sent by the sender.
- **gasPrice**
- **gasLimit**
- **to**: Recipient's address
- **value**: Amount of Wei Transferred from sender to recipient.
- **v,r,s**: Used to generate the signature that identifies the sender of the transaction.
- **init**: EVM code used to initialize the new contract account.
- **data**: Optional field that only exists for message calls.

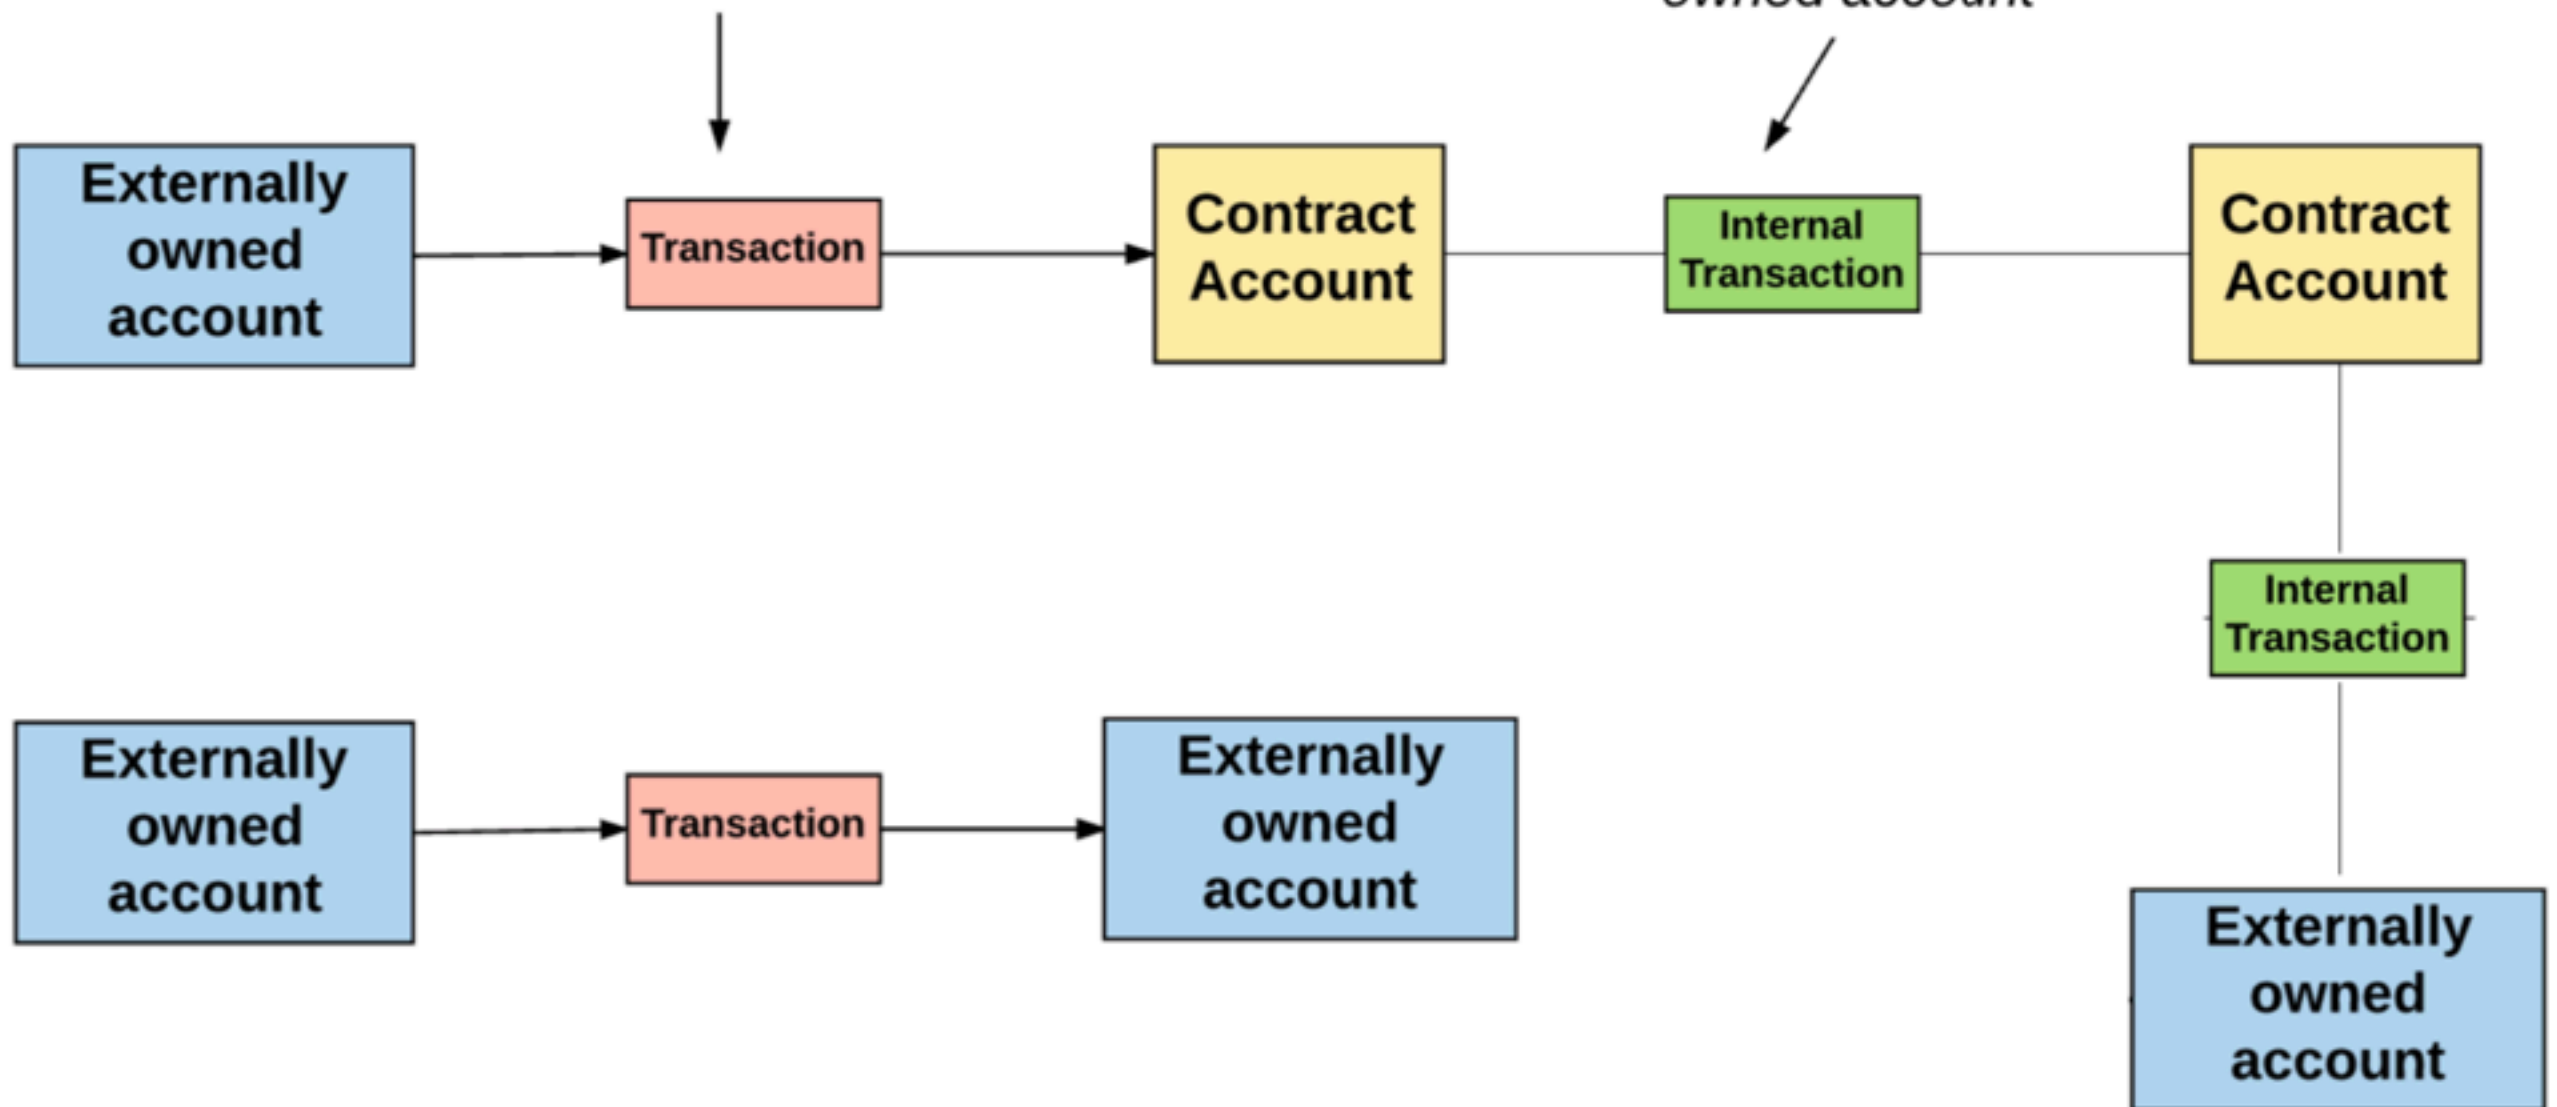
# Transaction Types

- External transactions
  - Generated by a user, serialized and sent to the Ethereum network, put on the blockchain (just like Bitcoin)
  - Includes contract creation, payment, contract calls
- Internal transactions
  - Contracts A can make a transaction for Contract B (aka, contract can “call a method” for Contract B)
  - These are not serialized and put on the blockchain!

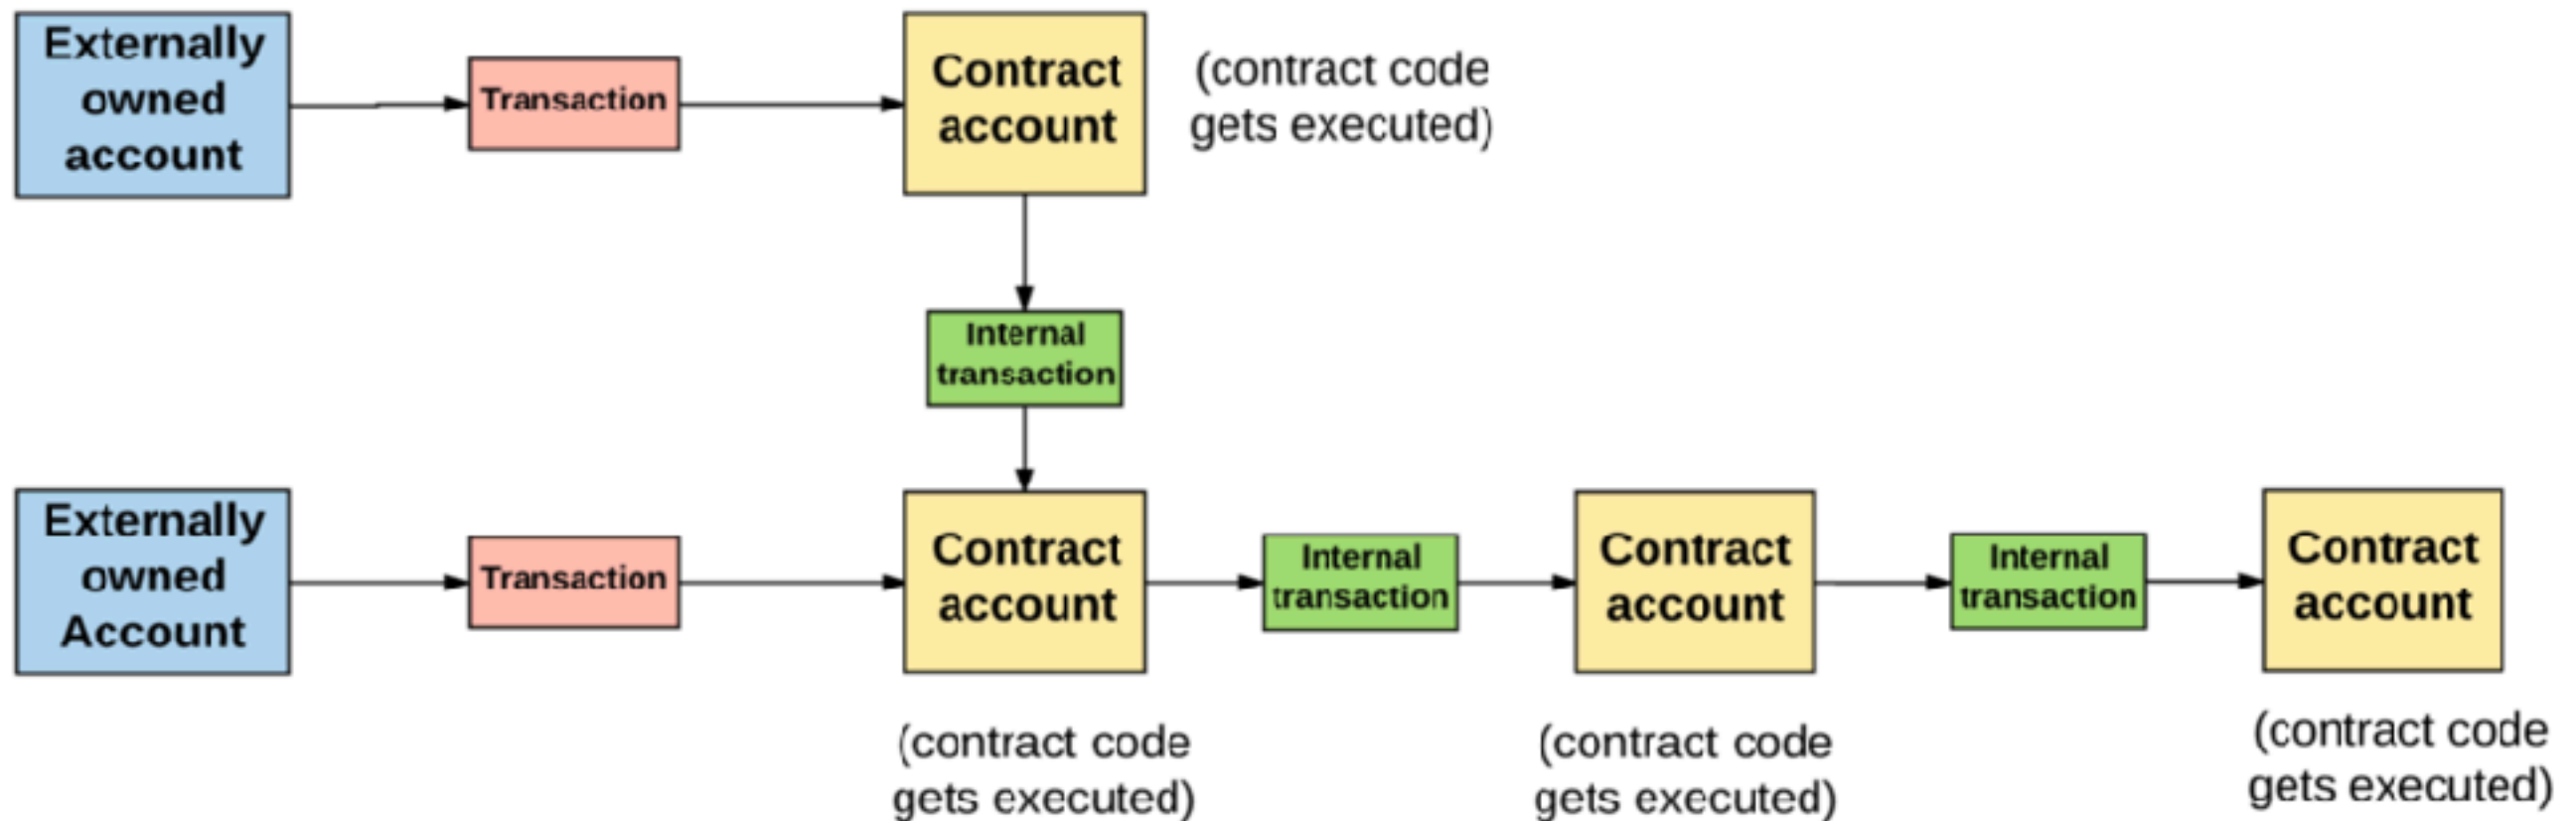
# Transaction Types

*Transaction set in motion by an externally owned account*

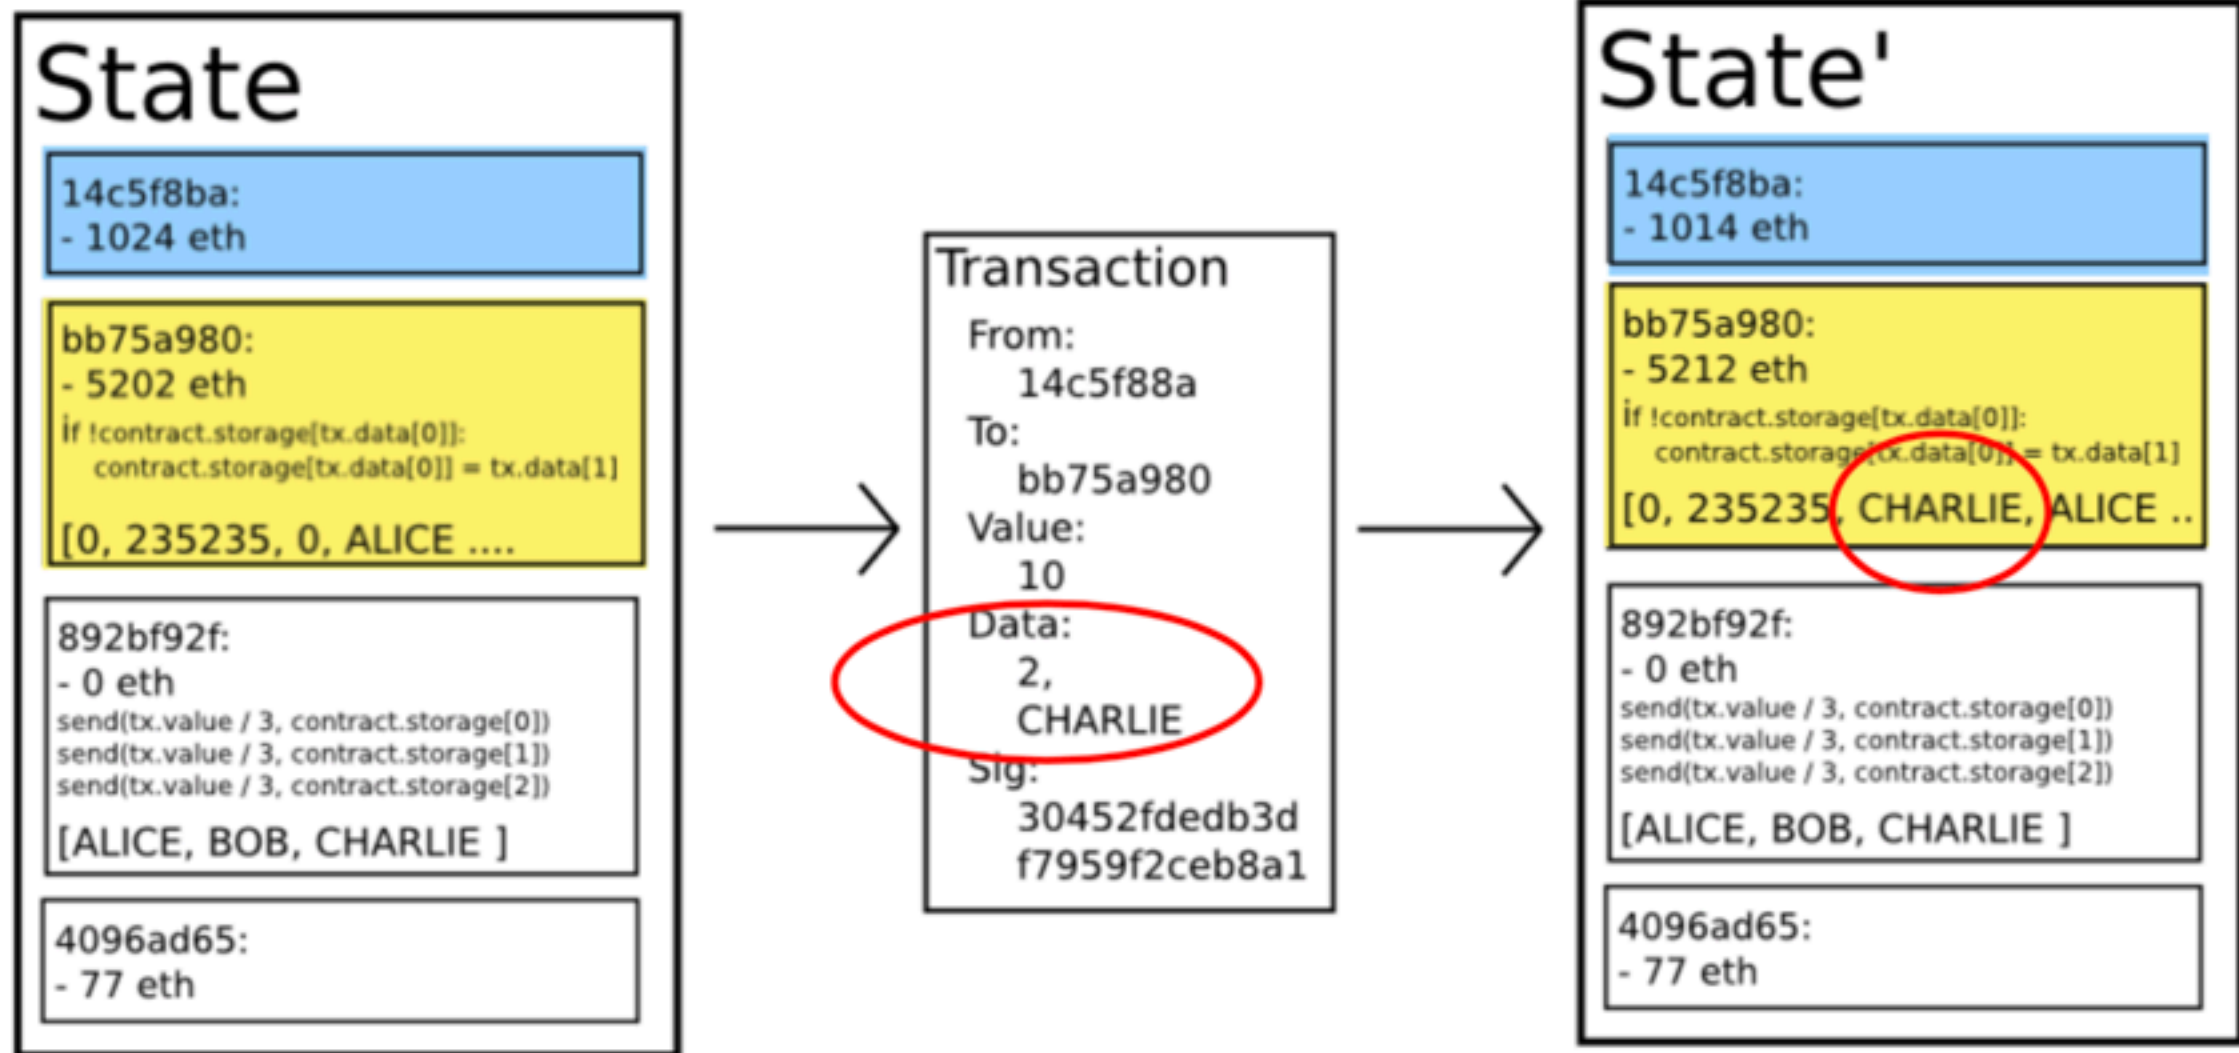
*Fired in response to the transaction set in motion by externally owned account*



# Transaction Flow



# Transaction Execution (Pay)



# Contracts

- Smart contracts are often written in a high-level object-oriented language (e.g., Solidity)
- The “contract” object has “methods”, which can be public or private (internal)
- Public methods can be called by anyone, private methods can only be called from other methods within that contract
- External transactions contain the contract address + the data (arguments) for a method call



- Smart contract  
object-oriented
- The “contract”  
public or private
- Public methods  
can only be called
- External transactions  
(arguments)

```

contract SimpleAuction {
    // Parameters of the auction. Times are either
    // absolute unix timestamps (seconds since 1970-01-01)
    // or time periods in seconds.
    address public beneficiary;
    uint public auctionEnd;

    // Current state of the auction.
    address public highestBidder;
    uint public highestBid;

    // Allowed withdrawals of previous bids
    mapping(address => uint) pendingReturns;

    // Set to true at the end, disallows any change
    bool ended;

    // Events that will be fired on changes.
    event HighestBidIncreased(address bidder, uint amount);
    event AuctionEnded(address winner, uint amount);

    // The following is a so-called natspec comment,
    // recognizable by the three slashes.
    // It will be shown when the user is asked to
    // confirm a transaction.

    /// Create a simple auction with `_biddingTime`
    /// seconds bidding time on behalf of the
    /// beneficiary address `_beneficiary`.
    constructor(
        uint _biddingTime,
        address _beneficiary
    ) public {
        beneficiary = _beneficiary;
        auctionEnd = now + _biddingTime;
    }

    /// Bid on the auction with the value sent
    /// together with this transaction.
    /// The value will only be refunded if the
    /// auction is not won.
    function bid() public payable {

```

methods

+ the data



# Contract Creation

- The same piece of code (“contract”) can be deployed multiple times, by different people
  - Contract “addresses” refer to a specific instance of a contract (combines contract code and a “nonce”)
  - Contracts are compiled into EVM byte code and sent to the network
  - To deploy the code, you send the code (as data) to the special Ethereum address (“0”)
    - (Contracts have a specialized opcode for this function...)

# EVM

To prevent cheating, the network works like Bitcoin:

Every single node in the network must also run the EVM machine instructions and inputs for each transaction in a received block, and only accepts the block if the EVM outputs (in the block) match their local computations.

Verification through repeated computing:

Each contract execution is “replicated” across the entire Ethereum network!

- What if that node cheats?

# How do we stop DoS?

- EVM code is “Turing complete”, e.g., has loops and may not halt
- Users pay fees to submit TXes to the network
  - (Pays for both computation and storage of data)
  - Contract computation/storage costs are denominated in “gas” (e.g., one hash call requires 30 “gas”)
  - Each transaction contains a “gas limit”, and a “gas price” price the sender will pay per unit of gas, denominated in ETH (currency)
    - Along with sufficient ETH to pay the fee

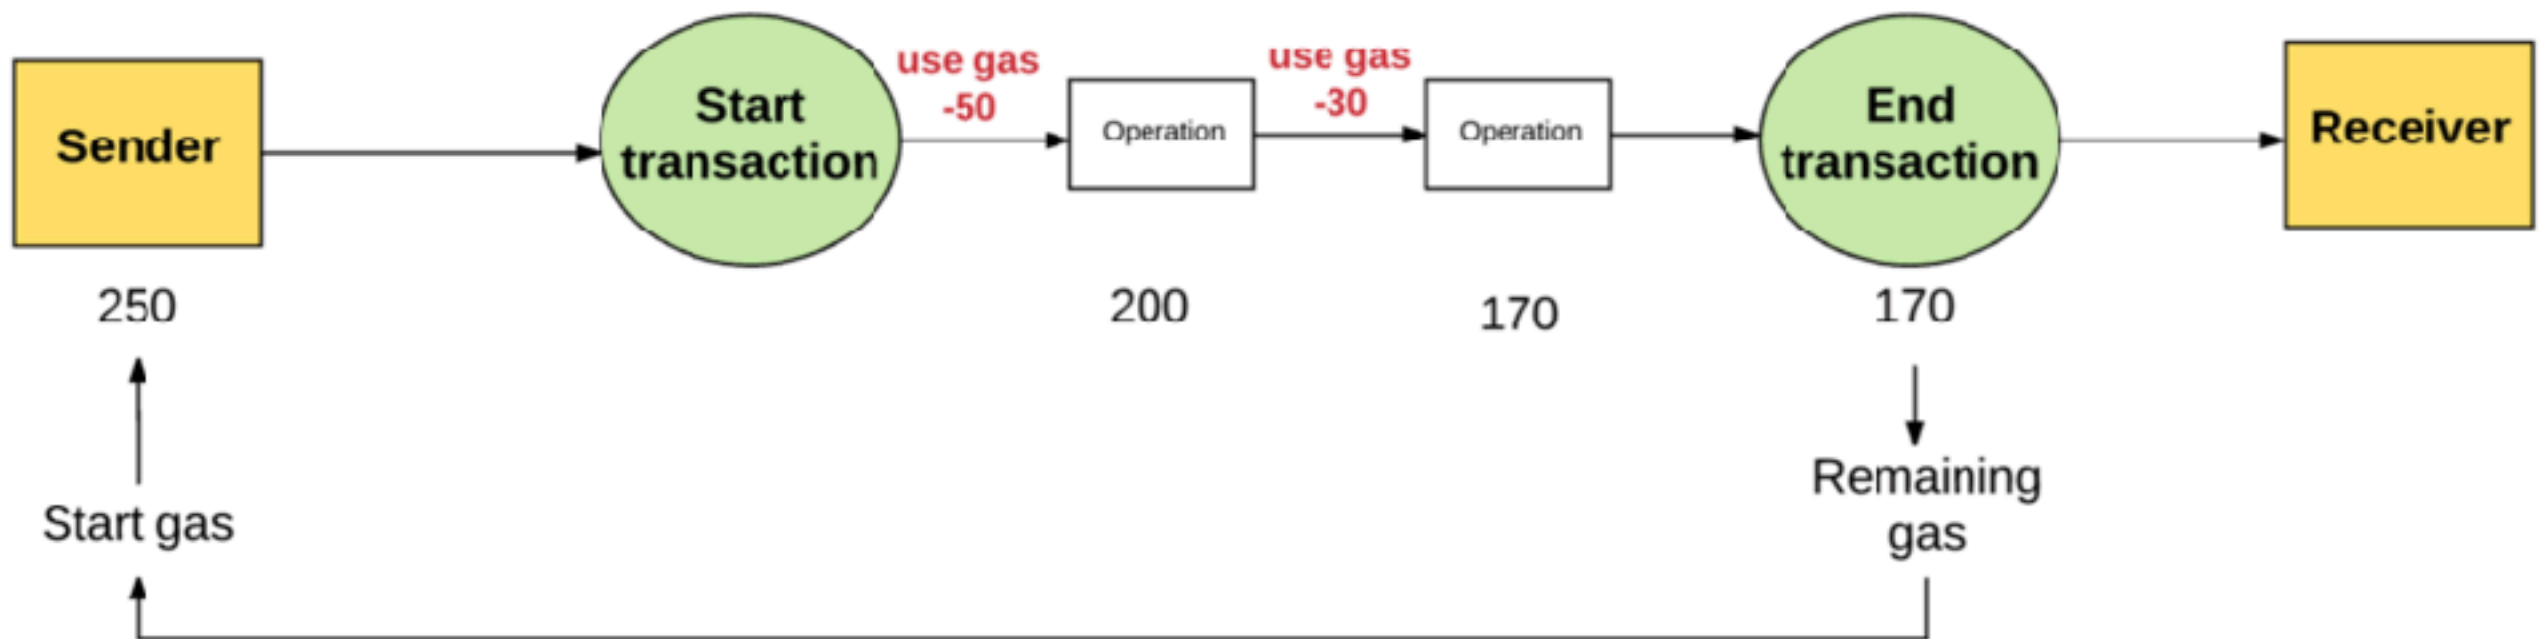
# Gas limits/price

- **Gas limit:** Max no. of computational steps the transaction is allowed.
- **Gas Price:** Max fee the sender is willing to pay per computation step.

Gas Limit		Gas Price		Max transaction fee
<b>50,000</b>	<b>x</b>	<b>20 gwei</b>	<b>=</b>	<b>0.001 Ether</b>


# Gas limits/price

The sender is refunded for any unused gas at the end of the transaction.



# Gas limits/price

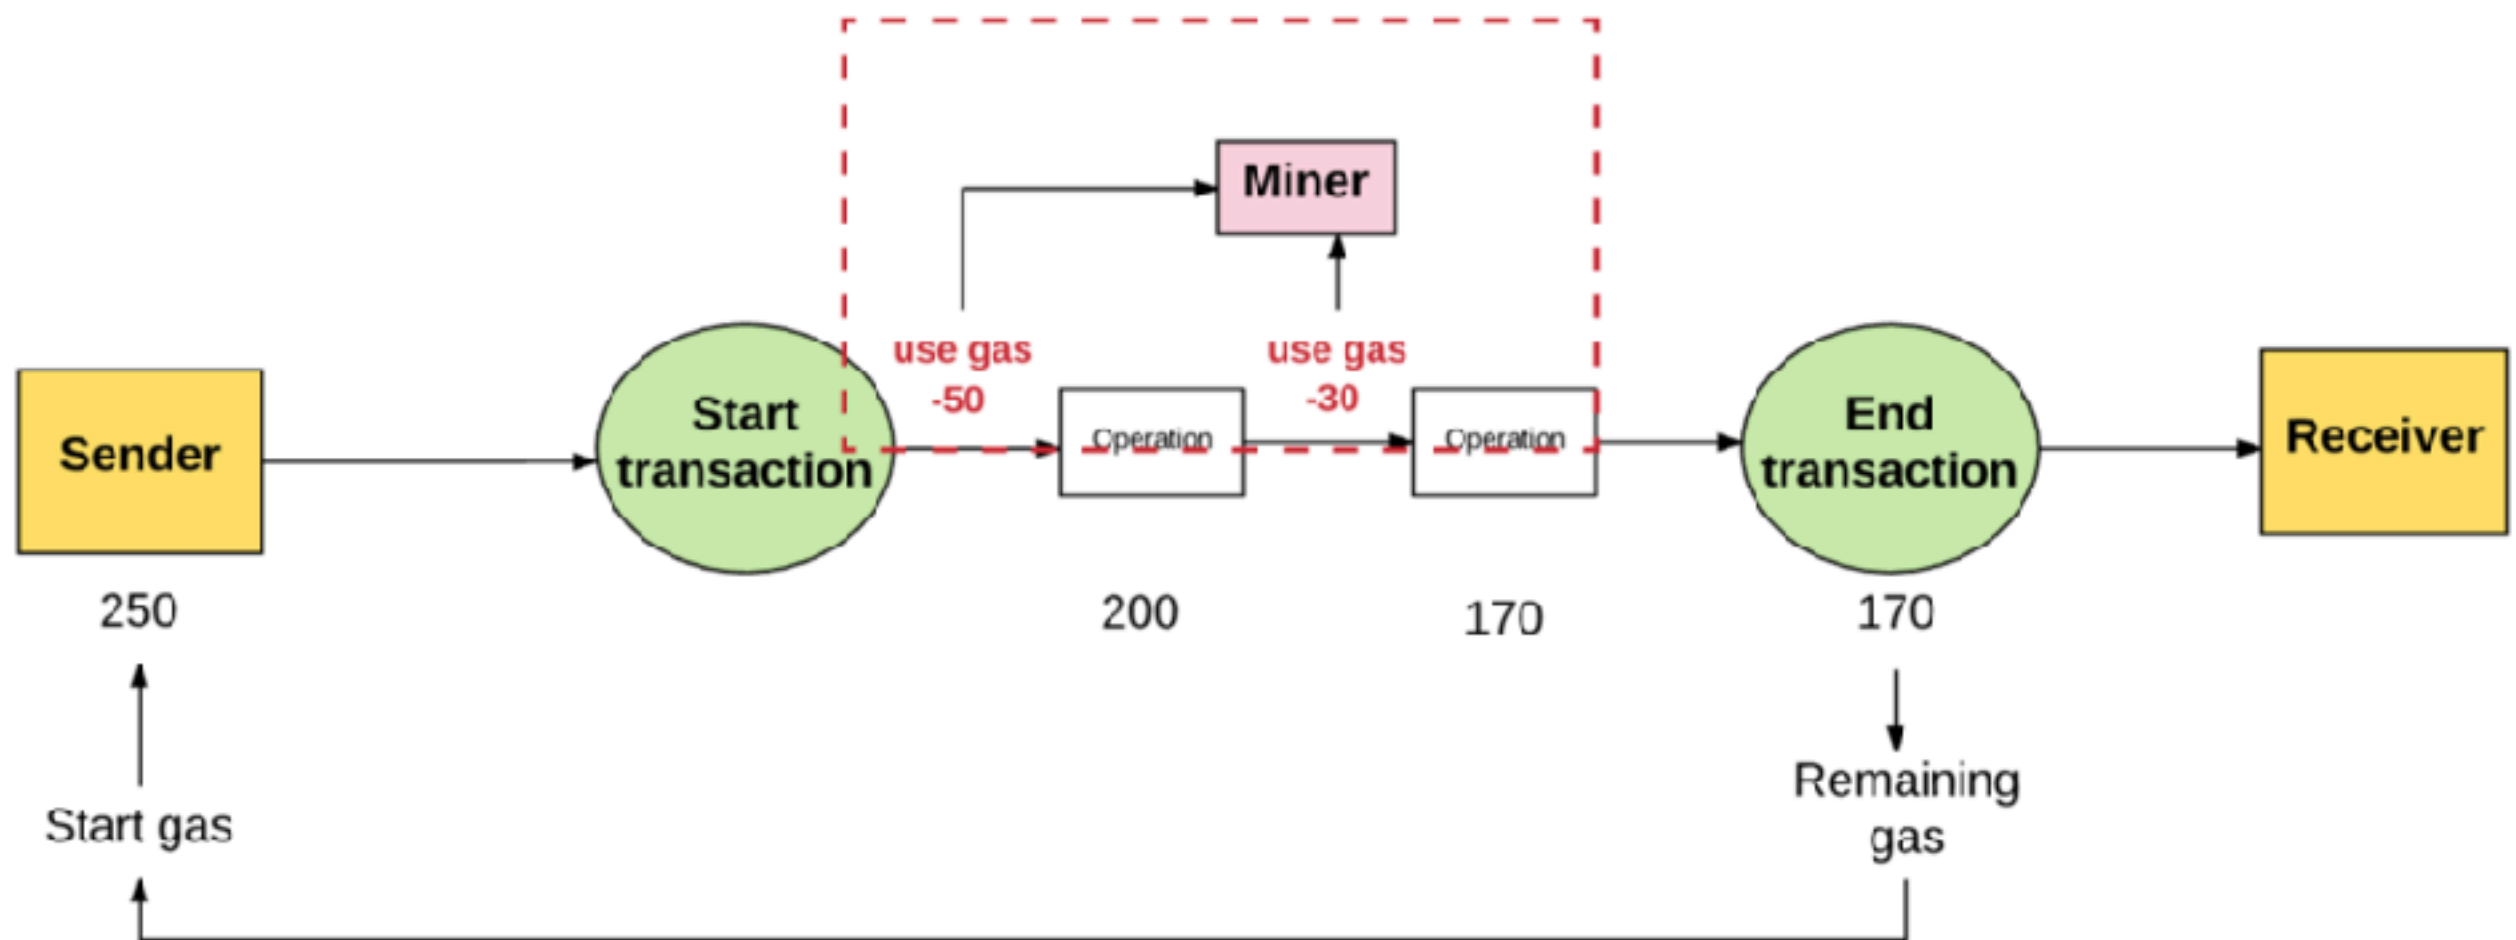
If sender does not provide the necessary gas to execute the transaction, the transaction runs “out of gas” and is considered invalid.



- The changes are reverted.
- None of the gas is refunded to the sender.

# Gas limits/price

All the money spent on gas by the sender is sent to the miner's address.



# Incentive problems?



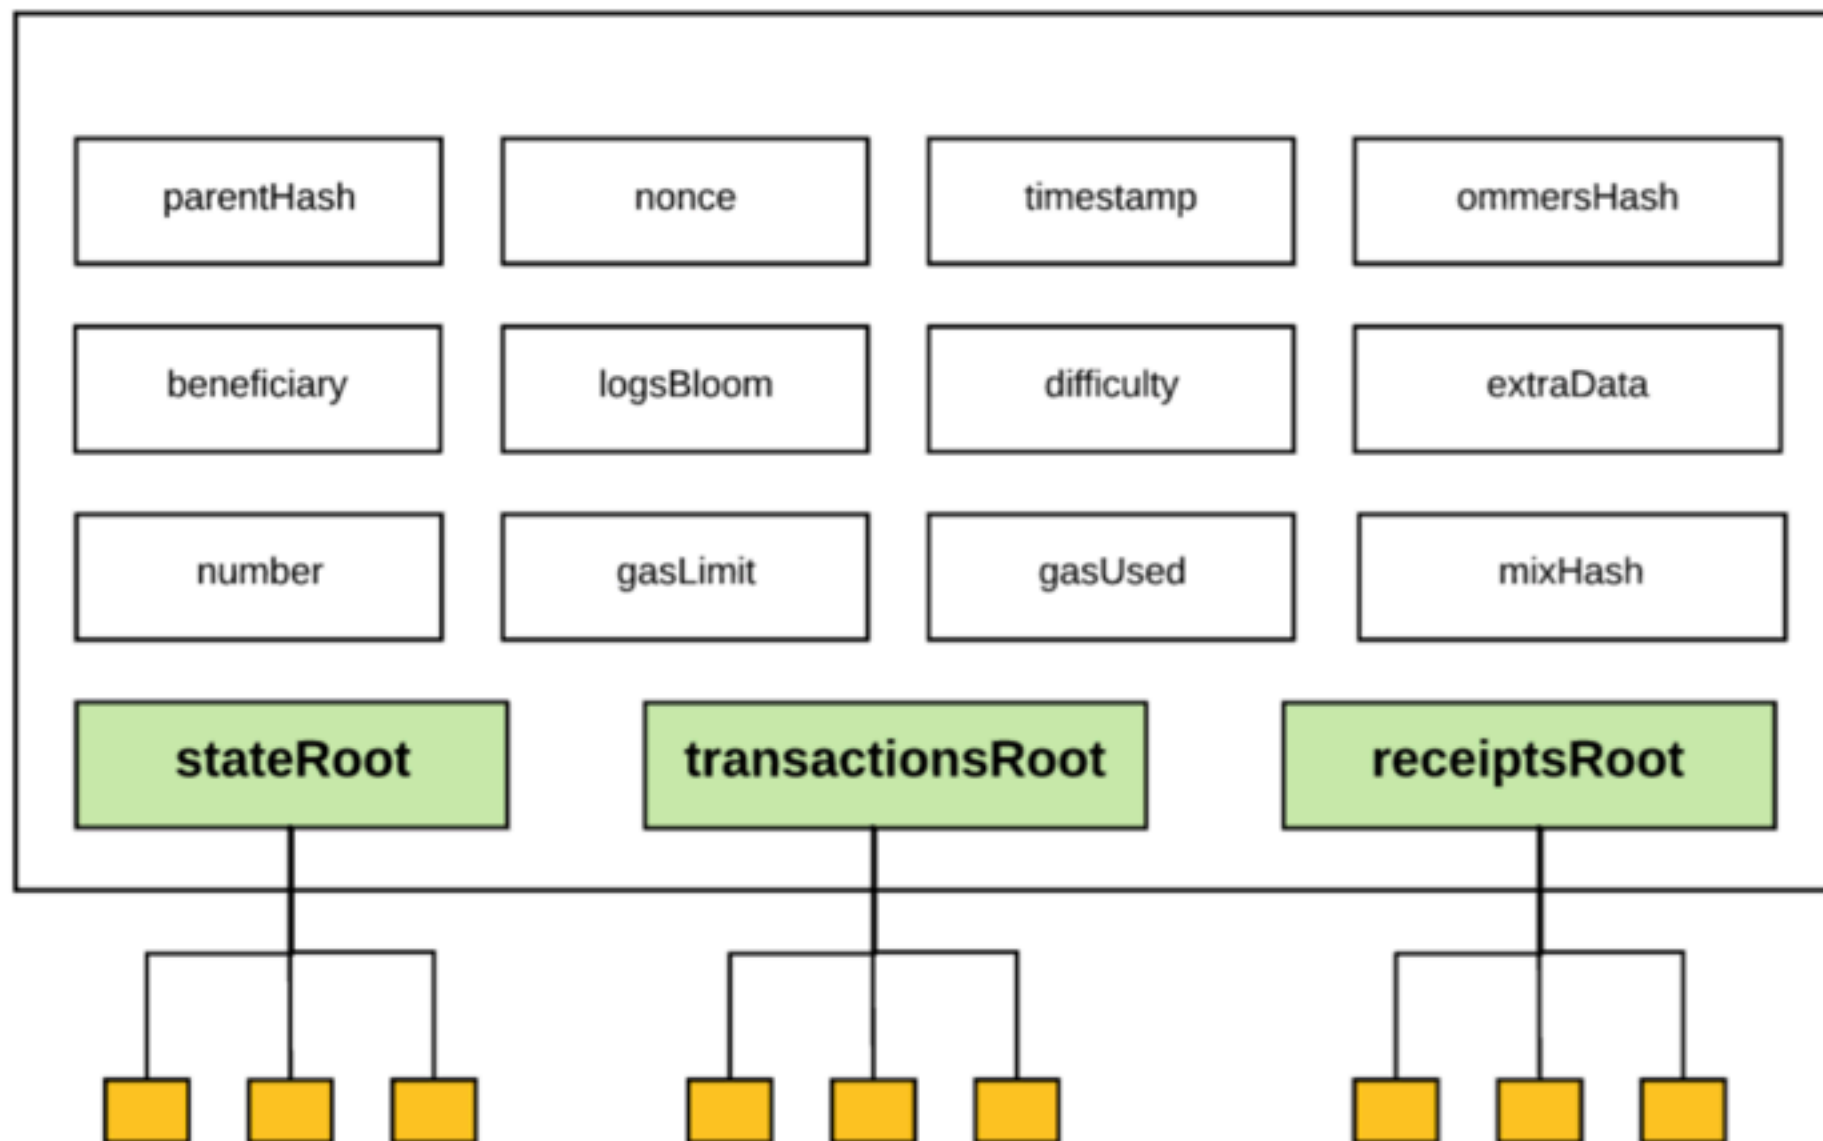
Does ETH need an internal  
currency?

# Where do internal TXes happen?

- Contracts can call public methods in other contracts, by authoring internal (“virtual”) transactions
- These transactions are not serialized and are not included on the blockchain
  - Instead, the miner formulates the internal transaction, executes the call locally, and serializes the result of the original call
  - The entire sequence of calls must be paid for within the calling transaction’s gas limit

# Merkle Trees

## Block header



## Binary Merkle Trees:

- Good data structure for authenticating information.
- Any edits/insertions/deletions are costly.

## Merkle Patricia Trees:

- New tree root can be quickly calculated after an insert, update edit or delete operation in  $O(\log n)$  time.
- Key-Value Pairs: Each value has a key associated with it.
- Key under which a value is stored is encoded into the path that you have to take down the tree.

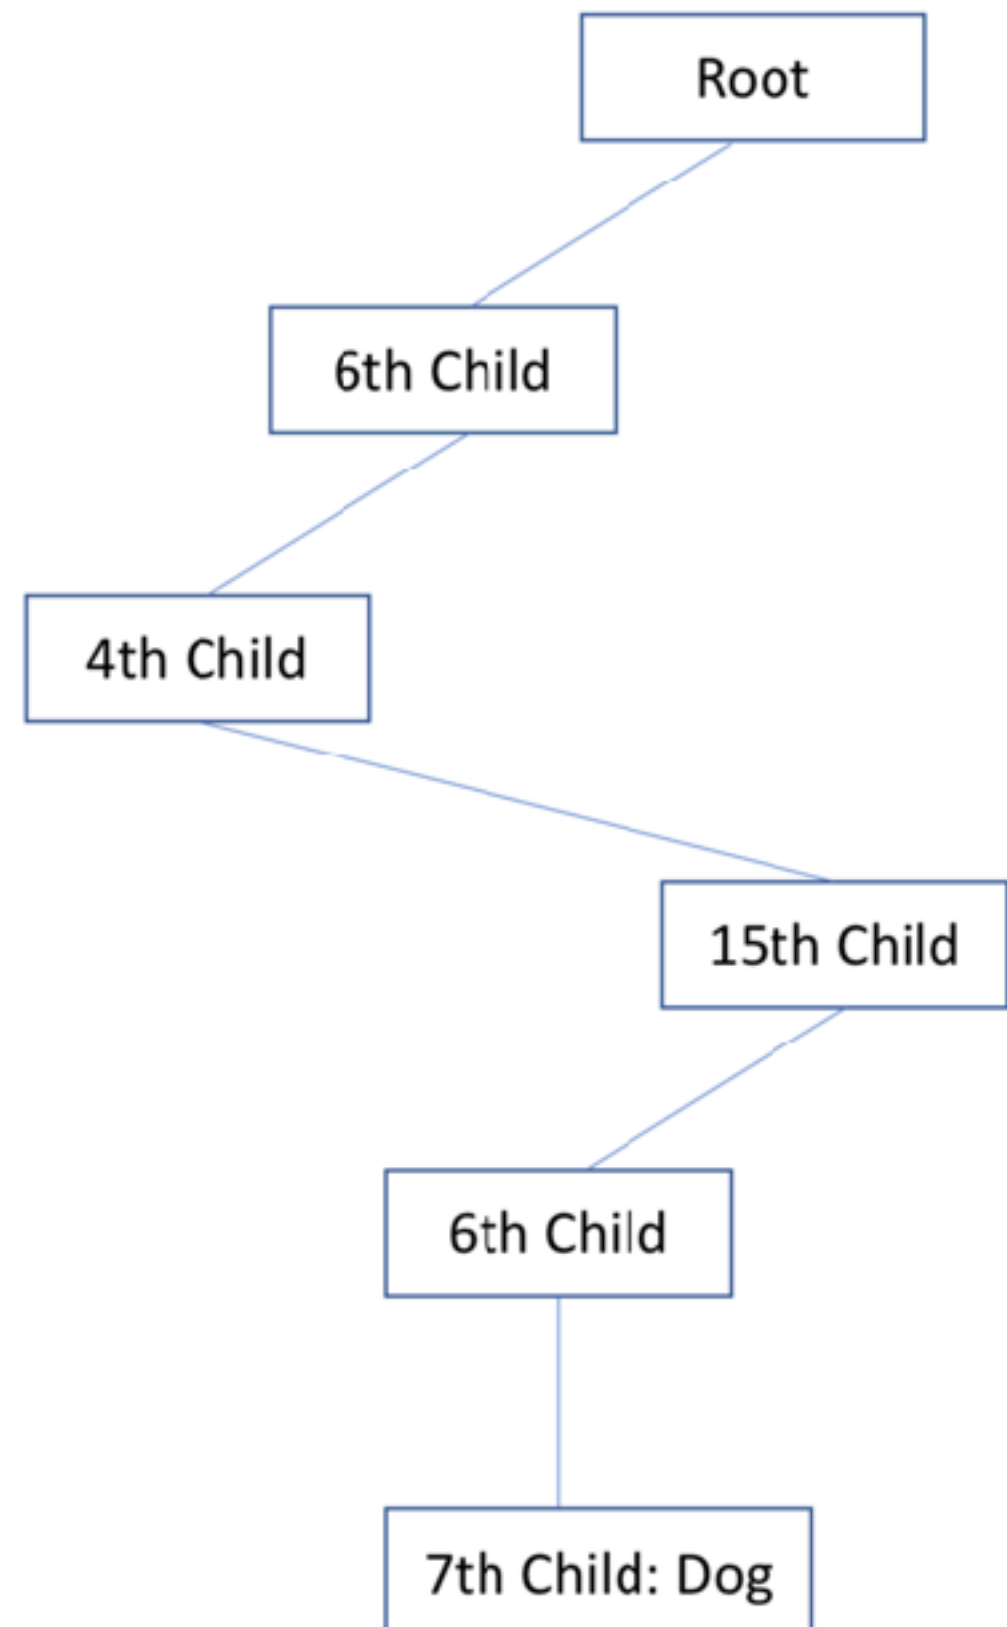
# Why Trees?

# Why Trees?

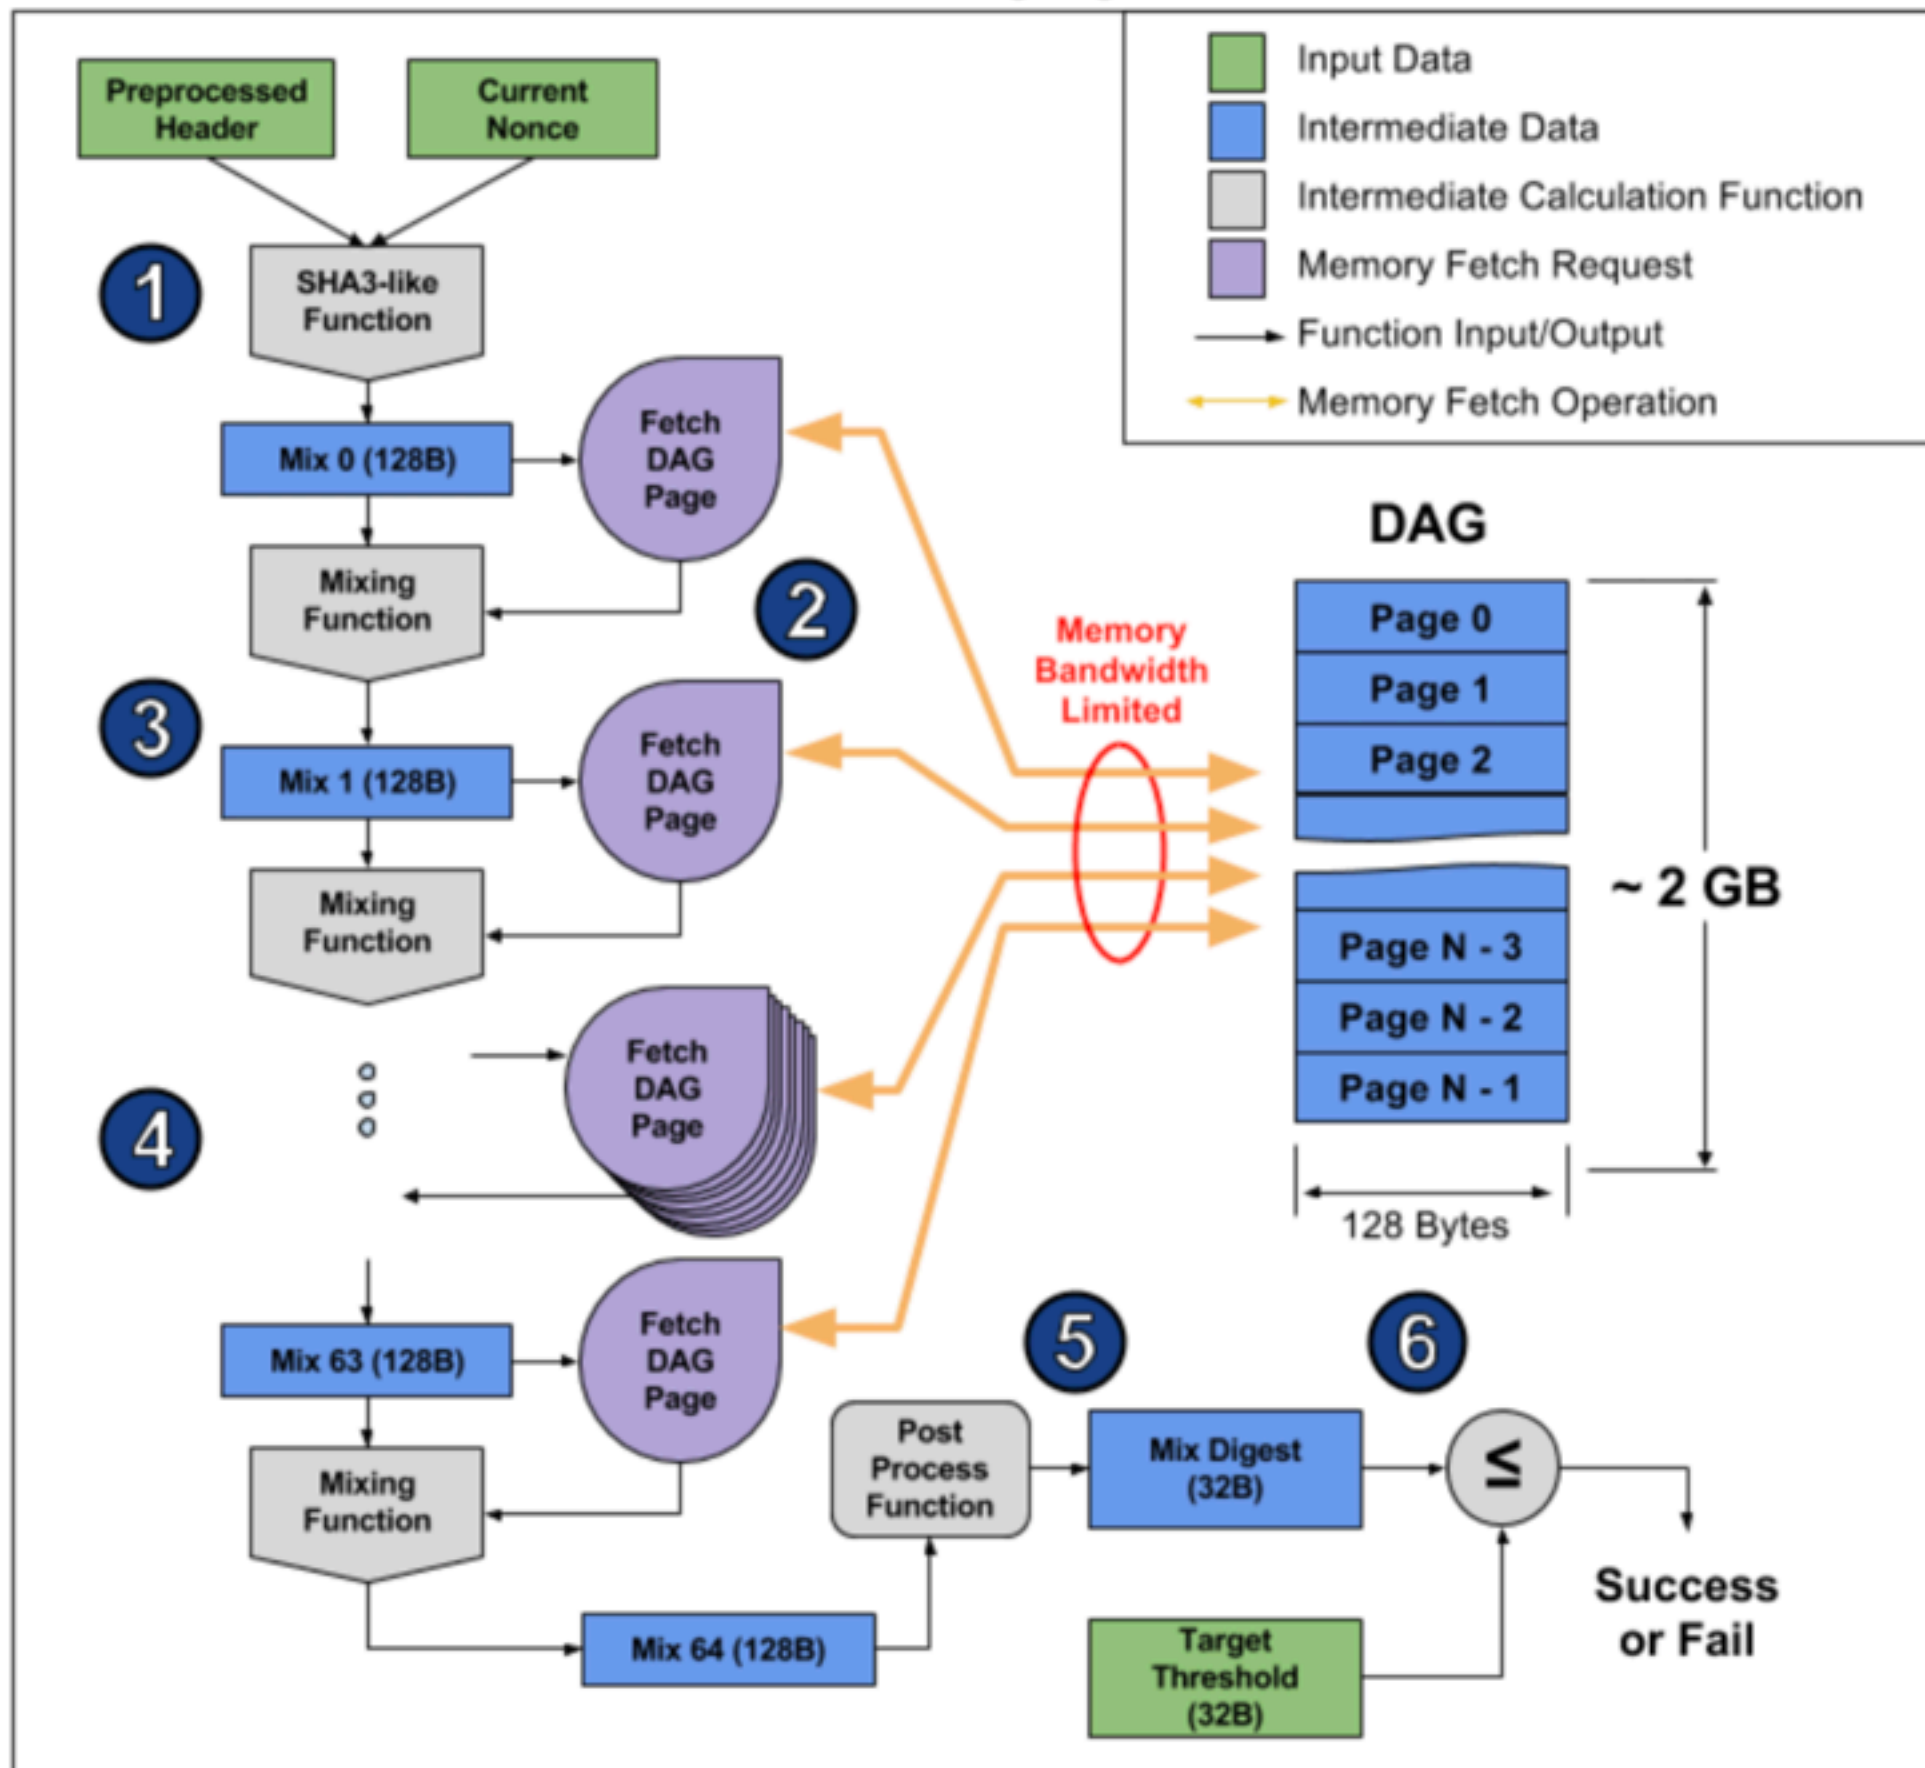
- A “light” client does not need to have the entire state of the world
- A server can construct a Merkle proof that data exists in the state tree
- Provided the state is encoded into the Merkle (Patricia) trees and we are given inclusion proofs, we can verify any block given just the previous block’s tree root
- Full nodes do need to keep current state data in order to work, but can “prune” (throw away) old states that are no longer valid

# Merkle Patricia Trees

- Each node has 16 children.
- eg: Hex(dog)= 6 4 6 15 6 7



# Ethash Hashing Algorithm





# Constructing the Chain

## Ommers/Uncles

- An ommer is a block whose parent is equal to the current block's parent's parent.
- Block times in Ethereum are around 15 sec. This is much lower than that in Bitcoin (10 min).
- This enables faster transaction. But there are more competing blocks, hence a higher number of **orphaned** blocks

# Constructing the Chain

## Ommers/Uncles

- An ommer is a block whose parent is equal to the current block's parent's parent.
- Block times in Ethereum are around 15 sec. This is much lower than that in Bitcoin (10 min).
- This enables faster transaction. But there are more competing blocks, hence a higher number of **orphaned** blocks
- The purpose of ommers is to help reward miners for including these orphaned blocks.
- The ommers that miners include must be within the sixth generation or smaller of the present block.

There are plans to replace Ethereum PoW with “proof of stake”, first through a “finality gadget” and later through a full PoS protocol.

This keeps getting delayed.

# Where do internal TXes happen?

- Contracts can call public methods in other contracts, by authoring internal (“virtual”) transactions
- These transactions are not serialized and are not included on the blockchain
  - Instead, the miner formulates the internal transaction, executes the call locally, and serializes the result of the original call
  - The entire sequence of calls must be paid for within the calling transaction’s gas limit

# Contract examples

- Simple “custom token” contract (ERC20)

```
1 // -----
2 // ERC Token Standard #20 Interface
3 // https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20.md
4 // -----
5 contract ERC20Interface {
6     function totalSupply() public view returns (uint);
7     function balanceOf(address tokenOwner) public view returns (uint balance);
8     function allowance(address tokenOwner, address spender) public view returns (uint remaining);
9     function transfer(address to, uint tokens) public returns (bool success);
10    function approve(address spender, uint tokens) public returns (bool success);
11    function transferFrom(address from, address to, uint tokens) public returns (bool success);
12
13    event Transfer(address indexed from, address indexed to, uint tokens);
14    event Approval(address indexed tokenOwner, address indexed spender, uint tokens);
15 }
```



```

1 contract TokenContractFragment {
2
3     // Balances for each account
4     mapping(address => uint256) balances;
5
6     // Owner of account approves the transfer of an amount to another account
7     mapping(address => mapping (address => uint256)) allowed;
8
9     // Get the token balance for account `tokenOwner`
10    function balanceOf(address tokenOwner) public constant returns (uint balance) {
11        return balances[tokenOwner];
12    }
13
14    // Transfer the balance from owner's account to another account
15    function transfer(address to, uint tokens) public returns (bool success) {
16        balances[msg.sender] = balances[msg.sender].sub(tokens);
17        balances[to] = balances[to].add(tokens);
18        Transfer(msg.sender, to, tokens);
19        return true;
20    }
21
22    // Send `tokens` amount of tokens from address `from` to address `to`
23    // The transferFrom method is used for a withdraw workflow, allowing contracts to send
24    // tokens on your behalf, for example to "deposit" to a contract address and/or to charge
25    // fees in sub-currencies; the command should fail unless the _from account has
26    // deliberately authorized the sender of the message via some mechanism; we propose
27    // these standardized APIs for approval:
28    function transferFrom(address from, address to, uint tokens) public returns (bool success) {
29        balances[from] = balances[from].sub(tokens);
30        allowed[from][msg.sender] = allowed[from][msg.sender].sub(tokens);
31        balances[to] = balances[to].add(tokens);
32        Transfer(from, to, tokens);
33        return true;
34    }
35
36    // Allow `spender` to withdraw from your account, multiple times, up to the `tokens` amount.
37    // If this function is called again it overwrites the current allowance with _value.
38    function approve(address spender, uint tokens) public returns (bool success) {
39        allowed[msg.sender][spender] = tokens;
40        Approval(msg.sender, spender, tokens);
41        return true;
42    }
43 }

```

# “NameCoin” in Ethereum

```
contract Namespace {  
  
    struct NameEntry {  
        address owner;  
        bytes32 value;  
    }  
  
    uint32 constant REGISTRATION_COST = 100;  
    uint32 constant UPDATE_COST = 10;  
    mapping(bytes32 => NameEntry) data;  
  
    function nameNew(bytes32 hash){  
        if (msg.value >= REGISTRATION_COST){  
            data[hash].owner = msg.sender;  
        }  
    }  
  
    function nameUpdate(bytes32 name, bytes32 newValue, address newOwner){  
        bytes32 hash = sha3(name);  
        if (data[hash].owner == msg.sender && msg.value >= UPDATE_COST){  
            data[hash].value = newValue;  
            if (newOwner != 0){  
                data[hash].owner = newOwner;  
            }  
        }  
    }  
  
    function nameLookup(bytes32 name){  
        return data[sha3(name)];  
    }  
  
}
```

Credit: Andrew Miller and Joe Bonneau for this code

# Multisig and filters

- Can create “filter” contracts that execute another contract and/or pay out money if complex conditions are satisfied
  - E.g., If k-out-of-N signers sign (in Bitcoin this is called “multisig”)
  - Verify that a certain number of blocks have elapsed
  - Check that another contract executed



# Prediction Markets

- Remember that contracts can “control” a balance (in ETH)
  - (They can control balances in e.g., ERC20 tokens as well)
- Can make payouts of ETH conditional on certain events — e.g., signed by a notary
  - Requires: method to “place bet”
  - Method to “claim bet” (verify sig/conditions), pay to an address
  - Relies on a centralized notary! See Augur....

# Decentralized Exchanges

- How do we build this?

# Frontrunning

# DAO disaster

- Decentralized Autonomous Organization
  - “Like a VC fund” but decentralized
  - Implementation: a contract that controls money, and directs its disbursement according to “shareholder votes”
  - Shareholders buy in, pool their ETH (sending to contract)
  - Then vote on investments, which are made together
  - Users can “split” a DAO

# “The DAO”

```
function splitDAO(
    uint _proposalID,
    address _newCurator
) noEther onlyTokenholders returns (bool _success) {

    ...
    // XXXXX Move ether and assign new Tokens. Notice how this is done first!
    uint fundsToBeMoved =
        (balances[msg.sender] * p.splitData[0].splitBalance) /
        p.splitData[0].totalSupply;
    if (p.splitData[0].newDAO.createTokenProxy.value(fundsToBeMoved)(msg.sender)
== false) // XXXXX This is the line the attacker wants to run more than once
        throw;

    ...
    // Burn DAO Tokens
    Transfer(msg.sender, 0, balances[msg.sender]);
    withdrawRewardFor(msg.sender); // be nice, and get his rewards
    // XXXXX Notice the preceding line is critically before the next few
    totalSupply -= balances[msg.sender]; // XXXXX AND THIS IS DONE LAST
    balances[msg.sender] = 0; // XXXXX AND THIS IS DONE LAST TOO
    paidOut[msg.sender] = 0;
    return true;
}
```

# How to upgrade a contract?

# How to upgrade a contract?

- Contracts are default immutable
- If there is no useful ongoing state, don't: just replace it
- If there is ongoing state (e.g., account balances) then:
  - Don't allow upgrades — and pray you got the code right
  - Call upgradeable/replaceable library code
  - Create a complex mechanism to transfer state from an old contract instance to a new contract instance

# Future of Ethereum

- Proof of stake
- Rollups
- Sharding